

DiVE: Diversifying View Recommendation for Visual Data Exploration

ABSTRACT

To support effective data exploration, there has been a growing interest in developing solutions that can automatically recommend data visualizations that reveal interesting and useful data-driven insights. In such solutions, a large number of possible data visualization views are generated and ranked according to some metric of importance (e.g., a deviation-based metric), then the top-k most important views are recommended. However, one drawback of that approach is that it often recommends similar views, leaving the data analyst with a limited amount of gained insights. To address that limitation, in this work we posit that employing diversification techniques in the process of view recommendation allows eliminating that redundancy and provides a good and concise coverage of the possible insights to be discovered. To that end, we propose a hybrid objective utility function, which captures both the importance, as well as the diversity of the insights revealed by the recommended views. While in principle, traditional diversification methods (e.g., Greedy Construction) provide plausible solutions under our proposed utility function, they suffer from a significantly high query processing cost. In particular, directly applying such methods leads to a “process-first-diversify-next” approach, in which all possible data visualization are generated first via executing a large number of aggregate queries. To address that challenge and minimize the incurred query processing cost, we propose an integrated scheme called *DiVE*, which efficiently selects the top-k recommended view based on our hybrid utility function. Specifically, *DiVE* leverages the properties of both the importance and diversity metrics to prune a large number of query executions without compromising the quality of recommendations. Our experimental evaluation on real datasets shows that DiVE can reduce the query processing cost by up to 40% compared to existing methods.

ACM Reference Format:

. 1997. DiVE: Diversifying View Recommendation for Visual Data Exploration. In *Proceedings of ACMconference (CIKM 2018)*. ACM, New York, NY, USA, Article 4, ?? pages. <https://doi.org/10.475/123.4>


Permission to make digital or hard copies of part or all of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for third-party components of this work must be honored. For all other uses, contact the owner/author(s).

CIKM 2018, October 2018,


© 2018 Copyright held by the owner/author(s).

ACM ISBN 123-4567-24-567/08/06...\$15.00

<https://doi.org/10.475/123.4>



(a) Visualization of the avg. oldpeak vs. chest pain types



(b) Visualization of the average age vs. chest pain types

Figure 1: Important vs. less important view.

1 INTRODUCTION

In the recent years, visualization recommendation systems have become an integral part of data exploration systems. The users who are interested in finding some meaningful insights in data have neither time nor patience to manually generate all possible data visualizations. In addition to time, the domain knowledge is another key factor when generating visualizations manually. However, with an exponential growth of available data in various domains, there has been an increase in the number of *Data Enthusiasts*, people with little domain knowledge and technical expertise, looking for interesting trends in data.

For instance, consider a Cleveland heart disease dataset¹, which describes patients with and without a heart disease. A data enthusiast might be interested in conducting some comparison between people with heart disease (disease) and people without heart disease (no disease). Without any prior

¹<http://archive.ics.uci.edu/ml/datasets/heart+Disease>




Figure 2: The visualization of maximum oldpeak vs. chest pain types

insights about data, she must manually specify different combinations of attributes, measures and aggregate functions before finally generating a visualization that reveals some interesting information about the dataset. The user effort and time spent in that process increases exponentially with increase in the number of attributes and measures. Hence, several data-driven visualization recommendation tools have been proposed to reduce the user effort and time during data exploration [? ? ?]. The main goal of those recommendation systems is to provide the user with the most important visualizations (top-k views), which are selected from all possible visualizations. The top-k views are selected based on the most important views in the dataset. The importance of a view is defined on the basis of particular criteria. One of the widely used criteria for importance is based on the deviation between the queried subset of data (target view) with the reference subset of data (reference view). The reference subset can be another subset of the dataset, the rest of the dataset, or the whole dataset. The intuition behind deviation based approach is that views that reveal substantially different trends from the reference views are likely to be of higher interest to the user[? ? ?].

Consider again the example of the heart disease dataset. Let the target subset be the data of people with heart disease and the reference subset be the data of people without heart disease. As shown in Figure ?? the average oldpeak (pressure of the ST segment) vs. chest pain types is more important view rather than the Figure ?? the average of age vs. chest pain types, due to the large deviation between target view (disease) and reference view (no disease) data. Figure ?? shows people with heart disease, especially who the chest pain types is asymptomatic tend to have much higher oldpeak rather than people without disease. To the contrary, the Figure ?? is potentially less important visualization compared to Figure ??, even there is a deviation between disease and no disease but the deviation is very small and if it is compared to Figure ??, it has lower deviation than Figure ??.

Figure ?? shows that there is no significant different in term of the average age of the people with four types of chest pain.

Although the deviation based visualization recommendation systems automatically provide users with the most important visualizations, it is likely that the views in the top-k set might be providing redundant information. For instance, Figure ?? provides the information which close to Figure ?? that the people with heart disease tend to have higher oldpeak values. Figure ?? has same attribute and attribute measure to Figure ??, the only difference is the aggregate function, where Figure ?? uses AVG and Figure ?? uses MAX. Since both views have a high deviation from the reference subset, both will appear in the top-k set. This leads to an important observation that using only importance as the selection criteria may deliver redundant recommended views, which leads to presents not optimal insights.

However, novelty and diversity are one of the fundamental characteristics of any effective recommendation systems [? ? ? ?]. Specifically, it is highly desirable that a visualization recommendation systems provides users with views that are both importance and also provide novel information that has not been revealed by the other views.

Towards designing an effective visualization recommendation systems that promotes both importance and novelty in recommended views, in this work, we propose an integrated approach called *DiVE*. In particular, *DiVE* aims to generate top-k visualizations that balance the tradeoff between importance and diversity. The main contributions of this paper are summarized as follows:

- We formulate the problem of evaluating recommended views that are both importance and diverse.
- We define a similarity measure to capture the distance between two visualizations.
- We present a hybrid objective function to balance the tradeoff between importance and diversity when ranking the visualizations.
- We propose the novel *DiVE* scheme, that employs various algorithms to evaluate the recommended visualizations based on the hybrid ranking/objective function.
- We present optimization techniques that leverage the hybrid objective function to substantially reduce the computational costs.
- We conduct an extensive experimental evaluation on real datasets, which compare the performance of various algorithms and illustrate the benefits achieved by *DiVE* both in terms of effectiveness and efficiency.

The rest of the paper is organized as follows: in Section II, we formulate the top-k diverse visualization problems and our related work; we present our proposed scheme *DiVE* in Section III; the experimental evaluation is reported in in Section IV and we conclude in Section V.

general notes:

1-use label to give each section and subsection a reference - important for roadmap and referring to sections.

2-For equations, no skips, no newlines and it should be one pair of \$ for each equation - some has \$ for each symbol, makes it hard to edit

3-add all missing references and follow the standard short format

2 PRELIMINARIES AND RELATED WORK

Several recent research efforts have been directed to the challenging task of recommending aggregate views that reveal interesting data-driven insights (e.g., []). As in previous work, we assume a similar model, in which a visual data exploration session starts with an analyst submitting a query Q on a multi-dimensional database D_B . Essentially, Q selects a subset D_Q from D_B by specifying a query predicate T . Hence, Q is simply defined as: $Q: \text{SELECT } * \text{ FROM } D_B \text{ WHERE } T$;

Ideally, the analysts would like to generate some aggregate views (e.g., bar charts or scatter plots) that unearth some valuable insights from the selected data subset D_Q . However, achieving that goal is only possible if the analyst knows exactly what to look for! That is, if they know the parameters, which specify some aggregate views that lead to those valuable insights (e.g., aggregate functions, grouping attributes, etc.). Meanwhile, such parameters only become clear in “hindsight” after spending long time exploring the underlying database. Hence, the goal of existing work, such as [? ? ? ? ?], is to *automatically* recommend such aggregate views.

To specify and recommend such views, as in previous work, we consider a multi-dimensional database D_B , which consists of a set of dimensional attributes \mathbb{A} and a set of measure attributes \mathbb{M} . Also, let \mathbb{F} be a set of possible aggregate functions over measure attributes, such as **COUNT**, **AVG**, **SUM**, **MIN** and **MAX**. Hence, specifying different combinations of dimension and measure attributes along with various aggregate functions, generates a set of possible views \mathbb{V} over the selected dataset D_Q . For instance, a possible aggregate view V_i is specified by a tuple $\langle A_i, M_i, F_i \rangle$, where $A_i \in \mathbb{A}$, $M_i \in \mathbb{M}$, and $F_i \in \mathbb{F}$, and it can be formally defined as: $V_i: \text{SELECT } A_i, F_i(M_i) \text{ FROM } D_B \text{ WHERE } T \text{ GROUP BY } A_i$;

Clearly, an analyst would be interested in those views that reveal interesting insights. However, manually looking for insights in each view $V_i \in \mathbb{V}$ is a labor-intensive and time-consuming process. For instance, consider again our example in the previous section. In that example, let D_B be the Cleveland Heart Disease table (i.e., **tb_heart_disease**) and the analyst is selecting the subset of patients with heart disease (i.e., $D_Q = \text{disease}$ subset). Hence, the number of views to explore is equal to: $|\mathbb{V}| = |\mathbb{A}| \times |\mathbb{M}| \times |\mathbb{F}|$, where $|\mathbb{F}|$ is the number of SQL aggregate functions, and $|\mathbb{A}|$ and $|\mathbb{M}|$ are the number of attribute and measures in **tb_heart_disease**, respectively. For that medium-dimensionality dataset, that value of $|\mathbb{V}|$ goes up to 180 views, which is clearly unfeasible for manual exploration. Such challenge motivated multiple research efforts that focused on automatic recommendation of views based on some metrics that capture the utility of a recommended view (e.g., [? ? ? ? ? ? ?]). *next sentences need to be more specific - one sentence for each of those works! The point is to show there is a space of recommendation methods and we are selecting the deviation-based one. Can come*

from your old related work section or from Humaira’s TKDE Some of those works focus on recommending visualizations to facilitate a particular user intent or task [? ? ?]. For example: explanations for a certain behavior, finding data anomalies or outliers and correlations among data attributes. Hence, the criteria for ranking the visualizations is driven by the user intent. However, in visual data exploration, often the intent of the user is not clear. Towards that end, data driven metrics are employed to capture the interestingness or importance of a recommended visualization.

Among the data driven metrics, recent case studies have shown that a *deviation-based* metric is effective in providing analysts with *important* visualizations that highlight some of the particular trends of the analyzed datasets [? ? ?].

In particular, the deviation-based metric measures the distance between $V_i(D_Q)$ and $V_i(D_R)$. That is, it measures the deviation between the aggregate view V_i generated from the subset data D_Q vs. that generated from a reference dataset D_R , where $V_i(D_Q)$ is denoted as *target* view, whereas $V_i(D_R)$ is denoted as *reference* view. That reference dataset could be the whole database (i.e., $D_R = D_B$) or a selected subset of the database. The premise underlying the deviation-based metric is that a view V_i that results in a high deviation is expected to reveal some important insights that are very particular to the subset D_Q and distinguish it from the patterns in D_R . In case, $D_R = D_B$, then the patterns extracted from D_Q are fundamentally different from the general ones manifested in the entire database D_B .

While recommending views based on their importance has been shown to reveal some interesting insight, it also suffers from the drawback of recommending similar and redundant views, which leaves the data analyst with a limited scope of the possible insights. *refer back to the intro example and reiterate that issue in one sentence* To address that limitation, in this work we posit that employing *diversification* techniques in the process of view recommendation allows eliminating that redundancy and provides a good and concise coverage of the possible insights to be discovered. In the next section, we discuss in details the formulation of both importance and diversity, and their impact on the view recommendation process.

check table-text consistency and make shorter - done also noticed that many symbols are often mentioned in the text without \$ especially K and S

3 DIVERSIFYING RECOMMENDED VISUALIZATIONS

short preamble

3.1 Content-Driven Deviation

As briefly described in the previous section, in this work we adopt a deviation-based metric to quantify the importance of an aggregate view [? ?]. Essentially, the deviation-based metric compares an aggregate view generated from the selected subset dataset D_Q (i.e., target view $V_i(D_Q)$) to the




Figure 3: Content vs. Context of views.

Table 1: Table of Symbols

Symbol	Description
k	no. of top recommended views
S	set of top- k recommended views
\mathbb{V}	set of all possible views
X	set of all candidate views
A	a dimensional attribute
M	a measure attribute
F	aggregate function
Q	a user query
D_B	a multi-dimensional database
D_Q	a target subset of D_B
D_R	a reference subset of D_B
V_i	a view query
$I(V_i)$	importance score of V_i
$I(S)$	importance score of views in S
$f(S, D)$	diversity score of views in S
$F(S)$	hybrid objective utility function value of S
$U(V_i)$	the utility score of each candidate view

same view if generated from a reference dataset D_R (i.e., reference view $V_i(D_R)$).

Clearly, the deviation between a target and a reference view is a *data-driven* metric. That is, it measures the deviation between the aggregate *result* of $V_i(D_Q)$ and that of $V_i(D_R)$. Consequently, and from a visualization point of view, that deviation is a *content-based* metric that captures the difference between the content of the visualization generated by $V_i(D_Q)$ vs. the visual content generated from $V_i(D_R)$. In the next, we formally describe the standard computation of that data-driven content-based metric, whereas the discussion of

its counterpart context-driven metric is deferred to the next section.

To calculate the content-based deviation, each target view $V_i(D_Q)$ is normalized into a *probability distribution* $P[V_i(D_Q)]$ and similarly, each reference view into $P[V_i(D_R)]$. In particular, consider an aggregate view $V_i = \langle A_i, M_i, F_i \rangle$. The result of that view can be represented as the set of tuples: $\langle (a_j, g_j), (a_j, g_j), \dots, (a_t, g_t) \rangle$, where t is the number of distinct values (i.e., groups) in attribute A_i , a_j is the j -th group in attribute A_i , and g_j is the aggregated value $F_i(M_i)$ for the group a_j . Hence, V is normalized by the sum of aggregate values $G = \sum_{j=1}^t g_j$, resulting in the probability distribution $P[V_i] = \langle \frac{g_1}{G}, \frac{g_2}{G}, \dots, \frac{g_t}{G} \rangle$.

Finally, the importance score of V_i is measured in terms of the distance between $P[V_i(D_Q)]$ and $P[V_i(D_R)]$, and is simply defined as:

$$I(V_i) = \text{dist}(P[V_i(D_Q)], P[V_i(D_R)]) \quad (1)$$

where $I(V_i)$ is the importance score of V_i and dist is a distance function. Similar to existing work (e.g., [?]), we adopt a Euclidian distance function, but other distance measures are also applicable (e.g., Earth Mover's distance, K-L divergence, etc.).

In current approaches for view recommendation, the importance value $I(V_i)$ of each possible view V_i is computed, and the k views with the highest deviation are recommended (i.e., *top-k*) (e.g., [?]). However, in this work, our goal is to ensure that recommended views provide a good coverage of possible insights, which is achieved by considering the context of the recommended views, which is described next.

3.2 Context-Driven Deviation

As mentioned above, recommending top-k views based only on their data content (i.e., content-driven deviation) often leads to a set of similar views. In order to provide full coverage of all possible interesting insights, in this work, we posit that achieving *diversity* within the set of recommended views is an essential quality measure. Diversity has been well known and widely used in recommendation systems for maximizing information gain and minimizing redundancy (e.g., [? ? ? ?]). At a high level, diversity essentially measures how different (i.e., diverse) are the individual data objects within a set.

Before discussing the details of diversity computation in Sec. ??, it is important to notice that central to that computation is some notion of distance measure between data objects. Existing work provides multiple metrics for measuring that distance between traditional data objects, such as web documents (e.g., [?]), database tuples (e.g., [?]), etc. However, our work in this paper is the first to consider diversity in the context of aggregate data visualizations. As such, a metric is needed to capture and quantify the (dis)similarity between the distinct features of different visualizations. Meanwhile, as each visualization is merely a data view generated by an SQL aggregate query, such metric naturally lends itself to considering the query underlying each view. That is, the query that has been executed to create the visualization. In turn, the distance between two visualizations is measured based on the distance between their underlying queries. Hence, in addition to the data-driven content-based deviation described above, here we also introduce a query-driven *context-based* deviation metric.

To measure the context-based deviation between two visualizations, we simply measure the distance between their underlying queries. Towards this, we extend on existing work in the area of query recommendation and refinement (e.g., [? ? ? ?]). In that work, the distance between two range queries q_1 and q_2 is mapped to that of measuring the edit distance needed to transform q_1 into q_2 , where the set of allowed transformation are: add, delete, or modify a predicate. In the context of our work, however, views are generated from aggregate queries without range predicates. In particular, a view is fully defined in terms of a combination of attribute, measure and an aggregate function. Hence, in addition to the content of a view V_i which is described by its probability distribution (i.e., $P(V_i)$) as defined in Sec ??), we also consider the context of the view $E(V_i)$, which is defined in terms of the query underlying V_i as: $E(V_i) = [A_i, M_i, F_i]$.

Such definition of view context leads to a special case of the existing work on query recommendation (e.g., [? ? ? ?]), in which the normalized distance between two queries is simply measured using the *Jaccard* similarity measure. Hence, the Jaccard similarity between two aggregate views V_i and V_j is measured as: $J(V_i, V_j) = \frac{|E(V_i) \cap E(V_j)|}{|E(V_i) \cup E(V_j)|}$

We note that the jaccard similarity assigns equal weights to each of the element in a set. Accordingly, when applied to aggregate views, then two views with same attribute and

different measure and aggregate function will have same similarity score as any other pair of views with same measure but different attribute and aggregate function. However, an analyst may consider two views with the same attribute A_i more similar than two views with same measure attribute M_i . To allow the analyst to specify such preference, each contextual component of a view is associated with a weight that specifies its impact on determining the (dis)similarity between views. Specifically, let w_i be the weight assigned to i^{th} element of set $C_{tx}(V_i)$, where $\sum_{i=1}^3 w_i = 1$. Then, the similarity between views V_i and V_j is measured as:

$$J(C_{tx}(V_i), C_{tx}(V_j)) = \frac{\sum_{i \in V_i \cap V_j} w_i}{\sum_{i \in V_i \cup V_j} w_i}$$

Consequently, the context-based deviation between V_i and V_j is calculated as:

$$D(C_{tx}(V_i), C_{tx}(V_j)) = 1 - J(C_{tx}(V_i), C_{tx}(V_j)) \quad (2)$$

something needs to be said about the figure! Maybe one sentence in each of the previous section, then simple discussion here

3.3 Problem Definition

In this section, we formally define our problem for recommending diversified interesting aggregate views. Towards this, we first define the metrics to measure the performance of our proposed visualization recommendation system in terms of: 1) the quality of recommended visualizations, and 2) the processing cost incurred in computing those visualizations.

3.3.1 Hybrid Objective Function. Our hybrid objective function is designed to consider both the importance and diversity of the recommended views. Particularly, it integrates two components: 1) the total importance score of set S and 2) the diversity score of S .

The importance score of the a S is calculated as the average value of the importance measure of each view in S , as given in Eq.?. Hence, the total importance score of S is defined as:

$$I(S) = \frac{\frac{1}{k} \sum_{i=1}^k I(V_i)}{I_u}, V_i \in S$$

where I_u is the upper bound on the importance score for an individual view. The value of I_u is used to normalize the average importance score for set S .

In order to measure the diversity of a set of objects, several diversity functions have been employed in the literature [? ? ?]. Among those, previous research has mostly focused on measuring diversity based on either the average or the minimum of the pairwise distances between the elements of a set [?]. In this work, we focus on the first of those variants (i.e., average), as it maximizes the coverage of S . Hence, given a distance metric $D(V_i, V_j)$, as given in equation ??, the diversity of a set S can be simply measured as follows:

$$f(S, D) = \frac{1}{k(k-1)} \sum_{i=1}^k \sum_{j>i}^k D(V_i, V_j), V_i, V_j \in S$$

Since the maximum context-based deviation between any two views in equation ?? is 1.0, then dividing the sum of

distances by $k(k-1)$ ensures that the diversity score of set S is normalized and bounded by 1.0.

Next, we define our proposed hybrid objective function that captures both the importance and diversity of the set of recommended views S . Specifically, for a set of views $S \subseteq V$, our hybrid objective function is formulated as the linear weighted combination of the importance score, $I(S)$ and diversity score $f(S, D)$, and is defined as:

$$F(S) = (1 - \lambda) I(S) + \lambda f(S, D) \quad (3)$$

where $0 \leq \lambda \leq 1$ is employed to control the preference given to each of the importance and diversity components. For instance, a higher value of λ results in a set of more diverse views, whereas a lower value of λ generates a set of the most important views, which is likely to exhibit redundancy in the recommended views.

Given the hybrid objective function, our goal is to find an optimum set of views S^* that maximizes the objective function $F(S)$, which is defined as follows:

DEFINITION 1. Recommending diversified important views : *Given a target subset D_Q and a reference subset D_R , the goal is to recommend a set $S \subseteq V$, where $|S| = k$, and V is the set of all possible target views, such that the overall hybrid objective $F(S)$ is maximized.*

Given the definition above, the quality of the recommended set of views is measured in terms of the value of the hybrid objective function $F(S)$.

$$S^* = \underset{\substack{S \subseteq V \\ |S|=k}}{\operatorname{argmax}} F(S) \quad (4)$$

3.3.2 Cost of Visualization Recommendation. Existing research has shown that recommending aggregate data visualizations based on data-driven content-based deviation is a computationally expensive task [? ? ?]. Moreover, integrating diversification to the view recommendation problem, as described above, further increases that computational cost. In particular, the incurred processing cost includes the following two components:

- (1) Query processing cost C_Q : measured in terms of the time needed to execute and compare all the queries underlying the set of target views as well as their corresponding reference views (i.e., content-based deviation).
- (2) View diversification cost C_D : measured in terms of the time needed to compute all the pairwise distances between each pair of target views (i.e., context-based deviation).

Consequently, the total cost C_T for recommending a set of views is simple defined as:

$$C_T = C_Q + C_D$$

In principle, traditional data diversification methods that consider both relevance and diversity can be directly applied in the context of our problem to maximize the overall utility function formulated in Eq.???. For instance, in the context

of recommending web search, such methods are designed to recommend a set of diversified objects (e.g., web documents) that are relevant to the user needs. However, in that setting, the relevance of an object is either given or simply computed. To the contrary, in our setting for view recommendation, the importance of a view is a computational expensive operation, which requires the execution of a target and reference view. As such, directly applying those methods leads to a “process-first-diversify-next” approach, in which all possible data visualization are generated first via executing a large number of aggregate queries. To address that challenge and minimize the incurred query processing cost, in the next section we propose an integrated scheme called *DiVE*, leverages the properties of both the importance and diversity to prune a large number of a large number of low-utility views without compromising the quality of recommendations, as described next.

However, in this work, we optimize the query execution time by reducing the search space of candidate views. Our proposed pruning scheme leverages the properties of importance and diversity metrics to prune a large number of low-utility views. The details of the pruning scheme are presented later in the paper.

In this section, we formally define the problem of recommending a set of k views that are interesting as well as diverse among themselves. In particular, given a target subset \mathcal{D}_Q and a reference subset \mathcal{D}_R , set of all possible views V , our objective is to select a set $S \subseteq V$ where $|S| = k$, such that the Interestingness score of the views in S and their mutual diversity is maximized.

Making visualization recommendations, particularly based on both data and context, is a computationally expensive task. Although hundreds of visualizations are possible for a given subset of data, only a small fraction of the visualizations are actually of interest and are included in the top- k set. Consequently, a significant fraction of the query processing cost is incurred on evaluating low-utility visualizations. Hence, in this work along with the quality of set S , we also aim to minimize the query processing cost of computing S . In particular, we apply effective pruning techniques to reduce the search space of visualizations.

Next, we define the metrics to measure the performance of our proposed visualization recommendation system in terms of quality of recommended visualizations and processing cost of computing those visualizations.

3.3.3 Hybrid Objective Function. Our hybrid objective function is designed to consider both the importance and diversity of the recommended views. Thus, we split the objective function into two components: 1) total importance score of set S and 2) the diversity score of S .

The importance score of the set S is calculated as the average value of the importance measure of each view in S , as given in equation ???. Formally, the total importance score of S is defined as:

$$I(S) = \frac{\frac{1}{k} \sum_{i=1}^k I(V_i)}{I_u}, V_i \in S$$

where I_u is the upper bound on the importance score for an individual view. The value of I_u is used to normalize the average importance score for set S . (i dont understand why this normalization is needed.)

In order to measure the diversity of a set of objects, several diversity functions have been employed in the literature [? ?]. Among those, previous research has mostly focused on measuring diversity based on either the average or the minimum of the pairwise distances between elements of set [?]. We focus on the first of those variants (i.e., average), as it considers all the views in S . Given a distance matrix $D(V_i, V_j)$ as given in equation ??, the diversity of a set S can be measured by a diversity function $f(S, D)$ that captures the dissimilarity between the views in S , defined as:

$$f(S, D) = \frac{1}{k(k-1)} \sum_{i=1}^k \sum_{j>i}^k D(V_i, V_j), V_i, V_j \in S$$

The maximum contextual distance between two views as in equation ?? is 1. Therefore, dividing the sum of distances between all views in S by $k(k-1)$ ensures that the diversity score of set S is normalized.

Next, we define our proposed hybrid objective function that captures both importance and diversity of the set of recommended views S . Specifically, for a set of views $S \subseteq V$ an objective function is formulated as the linear weighted combination of the importance score, $I(S)$ and diversity function $f(S, D)$ which is defined as:

$$F(S) = (1 - \lambda) \cdot I(S) + \lambda \cdot f(S, D) \quad (5)$$

where $0 \leq \lambda \leq 1$ is employed to control the contribution importance and diversity in the hybrid objective function. The higher values of λ result in a set of more diverse views whereas lower values of λ generate a set of the most important views that might be similar to each other. Given the hybrid objective function, our goal is to find an optimum set of views S^* that maximizes the objective function $F(S)$:

Hence, the quality of the recommended set of views is measured in terms of the value of the hybrid objective function $F(S)$.

$$S^* = \underset{\substack{S \subseteq V \\ |S|=k}}{\operatorname{argmax}} F(S) \quad (6)$$

3.3.4 Cost of Visualization Recommendation. As mentioned earlier, the visualization recommendation is a computationally expensive process [? ?]. Adding diversification to the view recommendation further increases the computational cost. Hence, our goal is to minimize the processing cost of selecting the top- k views that are both important and diverse. We split the processing cost of visualization recommendation into following two components:

- (1) Query Processing cost C_Q : measured in terms of query execution time to generate the content of each view.

- (2) Diversification cost C_{div} : measured in terms of time to compute the contextual distance between each pair of view.

Consequently, the total cost of recommending a set of top- k views can be defined as:

$$Cost(S) = C_Q(S) + C_{div}(S)$$

As the query execution cost includes both I/O and CPU cost, it dominates the overall cost of computing set S . Therefore, our goal in this work is to minimize the query execution time. Several approaches have been proposed in literature to optimize the query execution time i.e, query results caching[?], pruning techniques [?] shared computations, combining target and reference query, combination of multiple aggregates, combination of multiple GROUP BY, and parallel query and execution, [? ?].

However, in this work, we optimize the query execution time by reducing the search sapce of candidate views. Ou proposed pruning scheme leverages the properties of importance and diversity metrics to prune a large number of low-utility views. The details of the pruning scheme are presented later in the paper.

4 PROPOSED METHODS

As discussed in the introduction, the current view recommendations [? ? ?] generated solely on the basis of importance score suffer from the redundancy problem. The extreme solution to overcome the redundancy in the top- k views in the set S is to select views such that the diversity score of S is maximized.

This leads to two extreme baseline solutions, one based only on the importance score of the views which is called "**Linear-Importance**" and second based on only diversity score of the views which is called "**Greedy-Diversity**".

Instead of uses one of those two extrem, *DiVE* scheme employed hybrid function for evaluating top- k views which captures both the importance as well as diversity. Moreover, *DiVE* also equipped with pruning scheme which can prune a large number of queries without reducing the quality of the result. Our proposed *DiVE* scheme described below.

4.1 DiVE-Greedy Scheme

In order to capture both importance and diversity in the recommended top- k views, *DiVE* employs a greedy construction algorithm to iteratively select views that maximize the hybrid objective function $F(S)$ as presented in equation 3, and it is called as **DiVE-Greedy** scheme.

The details of the *DiVE-Greedy* scheme are given in Algorithm ?. The key ingredient of any Greedy algorithm for solving an optimization problem is the objective function itself that needs to be maximized. In particular, *DiVE-Greedy* initializes the set S with two most distant views. The distance between all the views is calculated using the distance function as given in equation 2. In each iteration, a new candidate view is selected from remaining views X and added to S . Thus, *DiVE-Greedy* assigns a utility score to each candidate

Algorithm 1: *DiVE Greedy*

Input: Set of views V and result set size k
Output: Result set $S \geq V$, $|S| = k$

```

1  $S \leftarrow [V_i, V_j]$  get two most distant views;
2  $X \leftarrow [V \setminus S]$ ;
3  $i \leftarrow \text{len}(S)$ ;
4 while  $i \neq k$  do
5   for  $j$  in set  $X$  do
6      $\max_v \leftarrow \text{argmax}_F(X[j], S)$ ;
7   end
8    $S.\text{add}(\max_v)$ ;
9    $X.\text{remove}(\max_v)$ ;
10   $i \leftarrow i + 1$ ;
11 end
12 return  $S$ 

```

view which is based on the hybrid objective function $F(S)$ as defined in equation ?? . The utility score of each candidate view V_i in X is computed as:

$$U(V_i) = (1 - \lambda) \cdot I(V_i) + \lambda \cdot \text{setDist}(V_i, S) \quad (7)$$

$$\text{Where } \text{setDist}(V_i, S) = \frac{1}{|S|} \sum_{V_j \in S} D(V_i, V_j)$$

Thus, the view with highest utility score in each iteration is selected and added to S .

DiVE-Greedy cost. The costs of Greedy Construction algorithm has two components which are the query execution cost C_Q that computing the importance score of view and the diversity cost C_D that computing set distance of each view from the views already in S . The complexity of query execution cost is $O(n)$ as the content of each view is generated only once. Meanwhile, the diversity cost C_D is $O(kn)$ where k is the size of subset of views S and n is the number of all possible views.

Traditional diversification method applied approach such "process-first-diversity-next" which leads to generating all possible views and computing the importance score in advance, then the diversity is executed next. Although, greedy algorithm is very efficient, however, as the number of attributes \mathbb{A} , measures \mathbb{M} and aggregate functions \mathbb{F} increase, the number of views that need to be generated grows exponentially. Hence, without a good strategy, *DiVE-Greedy* suffers from the high costs of query execution.

Therefore, to overcome this issue, *DiVE-Greedy* employs static pruning strategy as elaborated next.

4.2 DiVE-Greedy-Static Pruning

In order to recommend a small subset of views, the importance score I of all possible views need to be computed. However, only few views are eventually recommended to the analyst and rest of the views are discarded. Clearly, this approach would hinder the performance of the view recommendation systems for high dimensional datasets. Thus, motivated by the need to reduce the number of views that

Views in $V \setminus S$	$\min U'$	$\max U'$
V_1	0.23	0.65
V_2	0.19	0.21 ✕
V_3	0.25	0.85
V_4	0.15	0.70
V_5	0.20	0.23 ✕
V_6	0.21	0.24 ✕
V_7	0.22	0.91

Figure 4: Max-Min Pruning: All views which has $\max U'$ less than the maximum of $\min U'$ will be pruned

need to be generated, *DiVE-Greedy* employs a static pruning technique.

The proposed pruning technique is based on the observation that the utility score of each view $U(V_i)$ is a weighted sum of two different measures; 1) the importance score of view $I(V_i)$ and 2) diversity score of a view from S $\text{setDist}(V_i, S)$. The diversity score of a view requires only CPU computations and is faster operation. Whereas, computing the importance score of a view by comparing the target view to the reference view incurs high cost which is dominated by I/O cost of the query executions.

DiVE-Greedy applied Max-Min pruning method which has been presented in [?], by leveraging the diversity score of a view to decide whether a view query should be executed or not, and it is called **DiVE-Greedy-Static**.

Max-Min utilize the maximum bound of importance score I_u and minimum bound of importance score $I_0 = 0$. In each iteration, the distance $\text{setDist}(V_i, S)$ between all remaining views in X to the current set is calculated. Hence, the utility score of each views is computed using the real value of distance $\text{setDist}(V_i, S)$ and the importance score I_u and $I_0 = 0$. Thus, it produces two version of utility score which are $\min U'$ and $\max U'$ of each candidate views, which can be formally defined as:

$$\begin{aligned} \min U'(V_i) &= (1 - \lambda) \cdot I_0(V_i) + \lambda \cdot \text{setDist}(V_i, S) \\ \max U'(V_i) &= (1 - \lambda) \cdot I_u(V_i) + \lambda \cdot \text{setDist}(V_i, S) \end{aligned}$$

If $\max U'$ of the candidate view less than the maximum of $\min U'$, then this view definitely will be pruned, the detail example can be seen in Figure ?? . This approach generates same set of recommended views as generated by the *DiVE-Greedy* without pruning. This is due to the fact that in each iteration *DiVE-Greedy* selects the view with the highest utility score to be added to S .

This Max-Min pruning approach has been mentioned has a good performance in pruning [?]. However, in that literature, the dataset have a large of points which the value of each point is very diverse to others. To the contrary, this work is quite different. The context of view only has three dimention, while using three dimention and computing the diversity score of

the small set of view S , the utility score may not diverse as in the literature. Perhaps the performance of pruning is not same as well as in the literature.

Moreover, although Greedy algorithms have been shown to be efficient and provide good approximations to the optimal solutions [?], [?], [?]. However, since Greedy is constructive type algorithm, it constructs the set S by adding a new candidate view, there is no guarantee that the new view selected in each iteration is the best view for the objective function $F(S)$. It is because the view which has the highest utility score not necessary be the best one that improve the objective function $F(S)$ (e.g: local optimum).

To overcome this Greedy limitation, in this work, we also proposed another scheme which based on swap technique. The discussion of the swap technique is elaborated next.

Algorithm 2: DiVE Swap

Input: Set of views V and result set size k
Output: Result set $S \geq V$, $|S| = k$

```

1  $S \leftarrow$  Result set of only importance or only diversity;
2  $X \leftarrow [V \setminus S]$ ;
3  $F_{current} \leftarrow 0$ ;
4  $improve \leftarrow True$ ;
5 while  $improve = True$  do
6   for  $i$  in set  $X$  do
7      $S' \leftarrow S$ ;
8     for  $j$  in set  $S$  do
9       if  $F(S') < F(S \setminus S[j] \cup X[i])$  then
10         $S' \leftarrow S \setminus S[j] \cup X[i]$ ;
11      end
12    end
13    if  $F(S') > F(S)$  then
14       $S \leftarrow S'$ 
15    end
16  end
17  if  $F(S) > F_{current}$  then
18     $F_{current} \leftarrow F(S)$ ;
19     $improve \leftarrow True$ ;
20  else
21     $improve \leftarrow False$ ;
22  end
23 end
24 return  $S$ 

```

4.3 DiVE-Swap Scheme

Swap is local search type algorithm and it has been known and used to maximize diversity in the literature [? ?]. This algorithm starts with a complete initial set S , and try to achieve better result by interchanging the remaining views in X to the current set S . If a view in X is able to give better objective function value $F(S)$, then this view can be joined to the current set and one view in the current set that has the lowest contribution to the $F(S)$ will be removed. The

details of *DiVE-Swap* algorithm can be seen in Algorithm ??.

We proposed two types of Swaps which are: 1) **DiVE-iSwap**, the underlying behind this scheme is, it has the initial set from the result of Linear-Importance which is importance score maximized. 2) **DiVE-dSwap** which is quite similar to DiVE-iSwap, however, this scheme is initialized by results of Greedy-Diversity, which is diversity maximized. In each iteration, the candidate view is exchanged to the current set S till the $F(S)$ is maximized.

DiVE-Swap cost. The costs of Swap algorithm is also depend on the query execution time C_Q of all possible views and the diversity computation C_D . The query cost C_Q is executed only once but the cost is high due to it needs I/O cost. However, the complexity of diversity computation C_D is $O(k^2)$ and the number of distance computation depends on the number of iterations of the swap and the number of views in X . In the worst case, swap algorithm can perform $O(k^n)$ iterations. Without any pruning scheme, the cost of *DiVE-iSwap* is same as *DiVE-dSwap* due to those both schemes only different in the initial set.

However, in terms of pruning scheme, those schemes are quite different. *DiVE-iSwap* utilize the results of Linear-Importance as the initial set. Hence, this algorithm cannot escape from executing all queries due to Linear-Importance needs to execute all possible views to get the results. However, the second proposed swap algorithm, *DiVE-dSwap* is initialized by the result of Greedy-Diversity. This algorithm does not execute any query to generate the results. Therefore, leveraging the properties of diversity for pruning in *DiVE-dSwap* scheme is possible. The proposed static pruning of *DiVE-dSwap* is discussed below.

4.4 DiVE-dSwap Static Pruning

Similar to *DiVE-Greedy-Static*, we employ static pruning technique in this scheme which is by leveraging the diversity score of a candidate view to decide whether a view query should be executed or not. The details *DiVE-dSwap-Static* technique as follows:

- Since the initial set of *DiVE-dSwap* is the result of Greedy-Diversity, all query views in the initial set need to be executed in order to get the objective function $F(S)$ of the current set S .
- Before start an exchanging view from X to the current set S . All views in X is sorted based on $setDist(V_i, S)$. This is to confirm that exchanging process starts from the view that has highest score of diversity to the current set. This is called as "top-1" technique.
- In order to exchange the candidate view to the current set, the importance score must be known by executing the query of the candidate view. Instead of executing query view, the maximum bound of importance score is used to compute the utility score of each view as in *DiVE-Greedy-Static* technique. Hence, the result is not the actual utility score but $maxU'$, which defined as: $maxU'(V_i) = (1 - \lambda) \cdot I_u(V_i) + \lambda \cdot setDist(V_i, S)$.

- The exchanging process is started by comparing $F(S)$ of the current set to $F(S)$ of new set as can be seen in Algorithm ?? lines 9 - 10.
- If candidate views in X , while using importance score I_u cannot perform better in terms of improving the objective function $F(S)$ to the current set S , those views will be pruned.

This method is valid due to if the maximum score of importance I_u is used and that view cannot improve $F(S)$ of the current set, then there is no reason to execute the view query to get the importance score I . This pruning scheme is called **DiVE-dSwap-Static**.

All proposed pruning techniques including *DiVE-Greedy-Static* and *DiVE-dSwap-Static* are using static value I_u as the bound. Hence, the pruning performance of this technique will not optimal while the value of I_u is far away from the real maximum value of importance in database.

To overcome this issue, we proposed adaptive pruning scheme as described in the next section.

4.5 Adaptive Pruning Scheme

The idea behind pruning scheme is to minimize the query execution, which is by early prune low quality views. There are two main parameters in pruning scheme: 1) weight of λ and 2) the max bound value I_u .

The weight of λ is important due to it determines the contribution of importance score and diversity score. For instance, assume that an analyst wants to get set of views from view recommendation and she uses $\lambda = 0.7$. The λ value equal to 0.7 means that the diversity contribution to the utility score is 70% and the contribution of importance score will be 30 %, as it can be seen in equation ???. Thus, by setting the λ to higher value, the contribution of importance score will be lower and more queries are pruned. This λ value is determined by analyst, however, in this experiment 0.5 is used as the default.

Secondly, the maximum bound I_u has been utilized in *DiVE-Greedy-Static* and *DiVE-dSwap-Static* to decide whether the query view need to be executed or not. Only view that can improve the $F(S)$ of current set while using I_u will be executed otherwise those are pruned. However, one drawback using static bound I_u in pruning technique is that if the bound is far away from the maximum score of importance in the dataset, the pruning cannot work optimal.

To overcome this issue, instead of using static bound I_u , we proposed adaptive pruning scheme that automatically adapts the bound to the real maximum importance score in the dataset.

The idea behind the adaptive pruning scheme is by setting the bound to the maximum bound I_u as in static pruning for the first time, however, this bound is changed to the real value of importance score after its query view executed. However, the problem occurs when the executed query has a small importance score and it is far from the maximum importance score in the dataset. Thus, it brings the pruning out of control. Hence, DiVE needs the strategy to ensure

that the selected importance score is as close as possible to the maximum importance score in the dataset. One of the approach that can be used is by selecting sample views to be executed then get the maximum importance score of view from those sample . This brings us to the question of how many samples are needed in order to get a view that has a maximum score from the dataset.

There are several literatures have been mentioned related to the confidence interval and the number of samples in the normal distribution [cite]. However, the importance score of candidate views in X is not in normal distribution. The highest importance score I_u is equal to $\sqrt{2}$ whereas the lowest is 0, and it is long tail distribution. Hence, we adopt the sampling method from this [cite] as our data is not in normal distribution, it is called as prediction interval (PI) which is similar as a confidence interval in normal distribution. The relation between PI and the number of samples defined as in equation ??.

$$PI = \frac{(N-1)}{(N+1)}, \text{ where } N = \text{Number of samples} \quad (8)$$

When analyst uses $PI = 80\%$, it means there are 9 sample views executed, 85 %, 90%, 95%, 97%, and 99% means 12, 20, 40, 60, and 200 samples executed respectively.

The details of how our adaptive pruning scheme works is described as follows:

- Firstly, as in *DiVE-dSwap-Static* that all query view in the initial set are executed in order to get the objective function $F(S)$ of the current set S and all candidate views in X is sorted based on $setDist(V_i, S)$.
- $maxU'$ of each view is computed by utilize the maximum bound of importance score I_u , $maxU'(V_i) = (1 - \lambda) \cdot I_u(V_i) + \lambda \cdot setDist(V_i, S)$.
- All views that can improve $F(S)$ to the current set S are executed in order to get the actual value of importance.
- While the number of views executed based on the PI which determined by analyst. For instance, analyst may use $PI = 97\%$, hence, bound is changed while the sum of number of candidate views and th number of views in the initial set equal to 60. While the executed views equal to 60, the bound is replaced by the maximum importance score of executed views.

Adaptive pruning in Greedy is called *DiVE-Greedy-Adaptive* whereas in Swap is called *DiVE-dSwap-Adaptive*.

5 EXPERIMENTAL EVALUATION

Experiment Setup. This experiment running on Windows 10 64 bit, Intel Core i7-7700 CPU @ 3.60 GHz, RAM 16 GB. The experiment is developed using Python and run on Python version 3.6.3 with PostgreSQL as the database engine. Two real datasets are used, as follows:

- Heart Disease Dataset ², it has 14 attributes in total which are 9 attributes \mathbb{A} and 5 measure attributes \mathbb{M} ,

²<http://archive.ics.uci.edu/ml/datasets/heart+Disease>

Table 2: Parameters testbed in the experiments

Parameter	Range (default)
datasets	Heart disease , Flights
sample queries	10
diversity weight ratio	3(A) : 2(M) : 1(F)
tradeoff weight λ	0.0, 0.2, 0.4, 0.5 , 0.6, 0.8, 1.0
result set (size of k)	5 , 15, 25, 35
prediction interval %	80 , 85, 90, 95, 97 , 98

A= sex, cp, fbs, restecg, exang, slope, ca, thal, num; M= age, trestbps, chol, thalach, oldpeak. This dataset has small number of rows (299 rows).

- Airline (Flights) Dataset ³, it has 7 attributes, 4 measure attributes. The detail as follows: A= year, month, week, day, carrier, origin, destination; M= arrivaldelay, departurdelay, weatherdelay, distance. This dataset has 855,632 number of rows.

5.1 Effectiveness Evaluation

In order to test the performance of our proposed approach, we run experiments using two real datasets: flights dataset and heart disease dataset. For each dataset, we execute ten random queries. Afterwards, the average of all results is calculated as the final result. We examine the performance of our proposed schemes in term of the quality of the result (effectiveness) and the efficiency. All parameters used in this experiments are summarized in Table ?? and the detail results of the experiments is elaborated next.

The impact of parameter weight λ to the overall $F(S)$. One of the advantage *DiVE* scheme is users can decide what kind of the results that they want by changing the value of weight λ . The main function of λ is this weight can be used to tradeoff between importance and diversity, it is explained in equation ?? . Figure ?? shows the impact of λ to the overall objective values $F(S)$ by fixed k equal to 5, where $0 \leq \lambda \leq 1$. The figure shows that in the beginning when the λ is close to 0, means the role is dominated by the importance score, Linear-Importance has high overall objective function value $F(S)$ but it will decline gradually while increasing of λ . To the contrary, the Greedy-Diversity is the opposite, it has the high overall objective function value when the λ equal to 1, which means the diversity gets the full role. Hence, there is a crossover between Linear-Importance and Greedy-Diversity, this crossover shows the role balancing between importance score and diversity score. Furthermore, our proposed schemes have stable performance in all various λ value and its $F(S)$ value always better than Linear-Importance and Greedy-Diversity.

The impact of k to the overall objective function value $F(S)$. It has been explained the impact of λ to the overall objective function value $F(S)$. This section discuss the overall objective function value $F(S)$ by increasing k . Figure ?? shows the overall objective values $F(S)$ by using different

of k . As shown in those figures, the value of $F(S)$ decrease while increasing the value of k . There are two reasons why this moment happens: 1) the importance score of the set $I(S)$ is generated by sum of all importance score of views and it is divided by the number of views= k , increasing the number of k means increasing the divisor. 2) the diversity score of a set of views $f(S, D)$ also depends on the number of k . While k is increasing, the diversity score of the set will be decreased, due to the score of $f(S, D)$ is yielded from the sum of all distance of views divided by $k * (k - 1)$. Hence, the more views selected, the probability to get similar views will be high and the score of diversity will be declined.

The most important point in this results is our *DiVE* schemes always have better overall objective function value $F(S)$ compared to the two extreme baselines in various number of k .

6 CONCLUSIONS

In this paper, we proposed *DiVE* scheme which the main purposes are to evaluates and optimizes the results of visualization recommendation systems with respect to importance and diversity. The advantage of *DiVE* is that analyst can set their preferences by changing the parameter to tradeoff between importance and diversity to get result set. We also performed an experimental study and present the results which focus on effectiveness and efficiency of our approach on real datasets. We proposed *DiVE* scheme which based on Greedy and Swap approach, *DiVE-iSwap* have the best performance in recommending result views but it has the highest costs due to this scheme executing all possible view from the dataset, this scheme can be used for the analyst who only cares about the results without worrying execution time. However, to the analyst who care about execution time, we proposed *DiVE-dSwap-Adaptive* and *DiVE-Greedy-Adaptive*, those schemes are able to decrease costs significantly without reducing the quality of results.

ACKNOWLEDGMENTS

This research is supported by the Indonesia Endowment Fund for Education (LPDP) as scholarship provider from the Ministry of Finance, Republic of Indonesia.

REFERENCES

- [1] Charles L.A et al. Clarke. 2008. Novelty and diversity in information retrieval evaluation. *SIGIR* (2008).
- [2] M Drosou and E Pitoura. 2010. Search Result Diversification. *SIGMOD Record* (2010).
- [3] H Ehsan, M. Sharaf, and Panos K. Chrysanthos. 2016. MuVE: Efficient Multi-Objective View Recommendation for Visual Data Exploration. *ICDE* (2016).
- [4] Sean Kandel, Ravi Parikh, Andreas Paepcke, Joseph M Hellerstein, and Jeffrey Heer. 2012. Profiler: Integrated statistical analysis and visualization for data quality assessment. In *Proceedings of the International Working Conference on Advanced Visual Interfaces*. ACM, 547–554.
- [5] Verena Kantere. 2016. Query Similarity for Approximate Query Answering. In *DEXA*.
- [6] Verena Kantere, George Orfanoudakis, Anastasios Kementsietsidis, and Timos K. Sellis. 2015. Query Relaxation across Heterogeneous Data Sources. In *CIKM*.

³<http://stat-computing.org/dataexpo/2009/the-data.html>

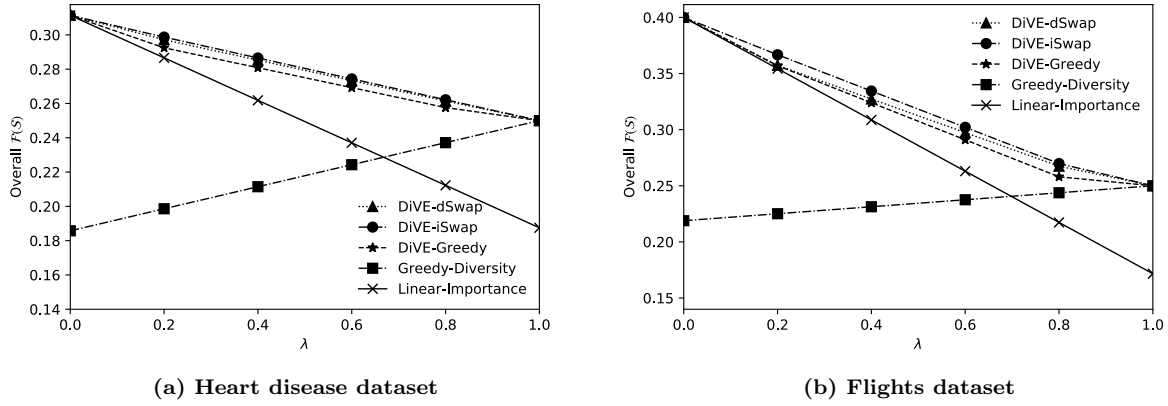


Figure 5: Impact of λ to overall objective function value $F(S)$ while $k = 5$ and running on three real datasets

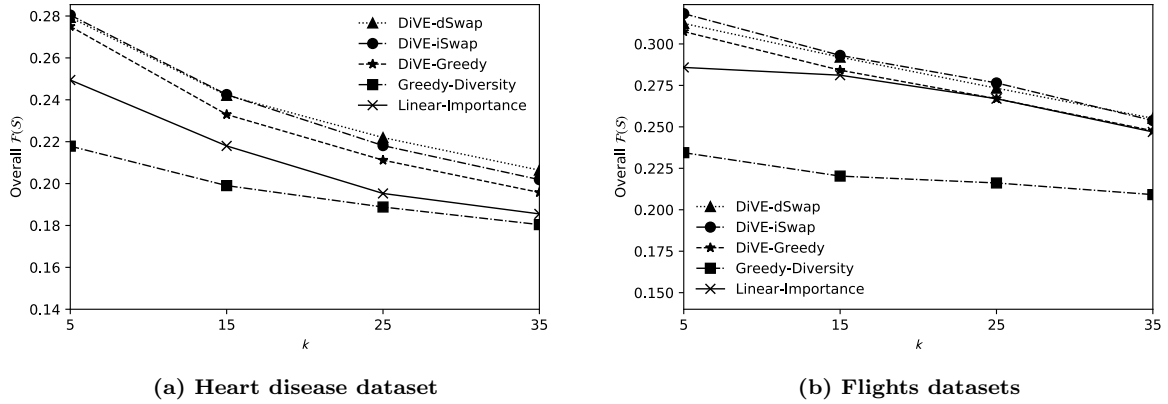


Figure 6: Overall objective function value $F(S)$ using different value of k and running on three real datasets

- [Alicia Key, Bill Howe, Daniel Perry, and Cecilia R. Aragon. 2012. VizDeck: self-organizing dashboards for visual analytics. *SIGMOD Conference* (2012).
- [Alicia Key, Bill Howe, Daniel Perry, and Cecilia R. Aragon. 2012. VizDeck: self-organizing dashboards for visual analytics. In *Proceedings of the ACM SIGMOD International Conference on Management of Data, SIGMOD 2012, Scottsdale, AZ, USA, May 20-24, 2012*. 681–684. <https://doi.org/10.1145/2213836.2213931>
- [Hina A. Khan and Mohamed A. Sharaf. 2015. Progressive diversification for column-based data exploration platforms. *ICDE* (2015).
- [Hina A Khan, Mohamed A Sharaf, and Abdullah Albarrak. 2014. DivIDE: Efficient Diversification for Interactive Data Exploration. *SSDBM* (2014).
- [Davood Rafiei, Krishna Bharat, and Anand Shukla. 2010. Diversifying web search results. *WWW* (2010).
- [Thibault Sellam and Martin L. Kersten. 2016. Fast, Explainable View Detection to Characterize Exploration Queries. In *Proceedings of the 28th International Conference on Scientific and Statistical Database Management, SSDBM 2016, Budapest, Hungary, July 18-20, 2016*. 20:1–20:12. <https://doi.org/10.1145/2949689.2949692>
- [Thibault Sellam and Martin L. Kersten. 2016. Ziggy: Characterizing Query Results for Data Explorers. *PVLDB* 9, 13 (2016), 1473–1476. <http://www.vldb.org/pvldb/vol9/p1473-sellam.pdf>
- [Jinwook Seo and Ben Shneiderman. 2006. Knowledge Discovery in High-Dimensional Data: Case Studies and a User Survey for the Rank-by-Feature Framework. *IEEE Trans. Vis. Comput. Graph.* 12, 3 (2006), 311–322. <https://doi.org/10.1109/TVCG.2006.50>
- [Barry Smyth and Paul McClave. 2001. Similarity vs . Diversity. (2001).
- [Quoc Trung Tran and Chee-Yong Chan. 2010. How to ConQueR why-not questions. In *SIGMOD*.
- [Manasi Vartak and Samuel Madden. 2014. S EE DB : Automatically Generating Query Visualizations. *VLDB* (2014).
- [M Vartak, S Rahman, S Madden, A Parameswaran, and N Polyzotis. 2015. SEEDB : Efficient Data-Driven Visualization Recommendations to Support Visual Analytics. *VLDB* 8 (2015).
- [Fernanda B. Viegas, Martin Wattenberg, Frank Van Ham, Jesse Kriss, and Matt McKeon. 2007. Many Eyes: A site for visualization at internet scale. *IEEE Transactions on Visualization and Computer Graphics* (2007), 1121–1128.
- [Marcos R. et al. Vieira. 2011. On query result diversification. *ICDE* (2011).
- [Eugene Wu, Leilani Battle, and Samuel R. Madden. 2014. The case for data visualization management systems. *VLDB Endowment* (2014).
- [Cong Yu, Laks Lakshmanan, and Sihem Amer-Yahia. 2009. It takes variety to make a world: diversification in recommender systems. *EDBT* (2009).

- Mi Zhang and Neil Hurley. 2008. Avoiding Monotony: Improving the Diversity of Recommendation Lists. *ACM Conference on Recommender Systems* (2008).