DiVE: Diversifiying View Recommendation for Visual Data Exploration

ABSTRACT

To support effective data exploration, there has been a growing interest in developing solutions that can automatically recommend data visualizations that reveal interesting and useful data-driven insights. In such solutions, a large number of possible data visualization views are generated and ranked according to some metric of importance (e.g., a deviation-based metric), then the top-k most important views are recommended. However, one drawback of that approach is that it often recommends similar views, leaving the data analyst with a limited amount of gained insights. To address that limitation, in this work we posit that employing diversification techniques in the process of view recommendation allows eliminating that redundancy and provides a good and concise coverage of the possible insights to be discovered. To that end, we propose a hybrid objective utility function, which captures both the importance, as well as the diversity of the insights revealed by the recommended views. While in principle, traditional diversification methods (e.g., Greedy Construction) provide plausible solutions under our proposed utility function, they suffer from a significantly high query processing cost. In particular, directly applying such methods leads to a "process-first-diversify-next" approach, in which all possible data visualization are generated first via executing a large number of aggregate queries. To address that challenge and minimize the incurred query processing cost, we propose an integrated scheme called *DiVE*, which efficiently selects the top-k recommended view based on our hybrid utility function. Specifically, *DiVE* leverages the properties of both the importance and diversity metrics to prune a large number of query executions without compromising the quality of recommendations. Our experimental evaluation on real datasets shows that DiVE can reduce the query processing cost by up to 40% compared to existing methods.

1 INTRODUCTION

In the recent years with an exponential growth of available data in various domains, there has been an increase in the number of people who trying to gain insights from the data by visual analysis (*data analyst*). Generaly, without any prior knowledge about the data, analyst must manually specify different combinations of attributes, measures and aggregate functions before finally generating a visualization that reveals some insights from the dataset.

Figure 1: Visualization of AVG**(oldpeak) vs. chest pain types**

However, manually looking for insights in each visualization is a labor-intensive and time-consuming process.

Such challenge motivated multiple research efforts that focused on automatic recommendation of visualizations based on some metrics that capture the utility of a recommended visualizations (e.g., [3, 5, 8, 10–12, 15–17]. Recent case studies have shown that *"a deviation-based metric"* to be effective in providing the most important visualization [15, 16]. The main goal of those works is to provide the most important visualizations (*top-k views*) to the analyst. In such solutions, a large number of possible views are generated and ranked according to that metric. This metric determines the deviation between the queried subset of data (*target view*) to the reference subset of data (*reference view*). The intuition behind deviation-based approach is that views that reveal substantially different trends from the reference view is judged as the important view.

+++ started here

Consider a data analyst trying to gain some insights into the *Cleveland heart disease* dataset[1]. Naturally, a first step in that exploratory analysis is to conduct some comparison between patients with heart disease (disease) and patients without heart disease (no disease). That is, disease is the target data subset for analysis and no disease is the reference data subset.

Since the analyzed data contains different dimensions (e.g.,) and different measures (e.g.,), it is a challenging task for the analyst to manually select the combination of dimensions and measures that reveal interesting insights. Hence, to automatically recommend interesting bar chart visualizations, different SQL aggregate functions are applied on the views generated from all the possible pairwise combinations of dimensions and measures, then the most important views are presented to the analyst. Figure 1 shows the top-1 important view according to the deviation-based metric [15, 16]. In particular, the figure shows that an aggregate view (i.e., bar chart) based on *average oldpeak* (i.e., pressure of the ST segment)

---

[1]http://archive.ics.uci.edu/ml/datasets/heart+Disease

Figure 2: Visualization of `MAX`**(oldpeak) vs. chest pain types**

vs. *chest pain types* exhibits a large deviation between the target view (`disease`) and reference view (`no disease`).

While recommending views based on their importance has been shown to reveal some interesting insights [3, 15, 16], such approach still suffers from a "tunnel vision", where it often recommends similar and redundant views. For instance, Figure 2 shows the second top recommended view for the analysis described above. Comparing Figures 2 and 1, it is easy to see that both visualizations are based on the same dimension (i.e., *chest pain types*) and the same measure (i.e., *oldpeak*), and the only difference between them is the aggregate function (i.e., `MAX` vs `AVG`). Despite that similarity between the two views, they are still both recommended to the analyst due to the high deviation between the target and reference views.

+++ stopped here

Although the deviation-based visualization recommendation automatically provide the top-k most important views, however, it often recommends similar views and leaving the analyst with a limited amount of gained insights. For instance, Figure 2 provides similar insights to Figure 1, both figures show that people with heart disease tend to have higher oldpeak values. Figure 2 is generated using *chest pain types* as the attribute, *oldpeak* as the measure and `MAX` as the aggregate function, whereas Figure 1, uses same attribute and measure but uses `AVG` as the aggregate function. As shown in both figures, since those two views have a high deviation from the reference view, both views are considered as the important views and those will appear in the top-k set. This leads to an important observation that using "only importance" as the only criterion (e.g. a deviation-based metric) is that often deliver redundant recommended views, which leads to presents limited insights of results.

To address that limitation, in this work we posit that employing diversification techniques in the process of view recommendation allows eliminating that redundancy and provides concise coverage of the possible insights to be discovered. In fact, novelty and diversity are one of the fundamental characteristics of any effective recommendation systems [1, 9, 20, 21]. Specifically, it is highly desirable that a view recommendation can recommend the top-k views that are both importance and also provide new insights that has not been revealed by the other views.

Towards designing an effective view recommendation that promotes both importance and diversity in recommended views, in this work, we propose an integrated approach called *DiVE*. In particular, DiVE aims to generate top-k views that balance the tradeoff between importance and diversity. The main contributions of this paper are summarized as follows:

- We formulate the problem of evaluating recommended views that are both importance and diverse (**Section 3**).
- We define a similarity measure to capture the distance between two visualizations (**Section 3**).
- We present a hybrid objective function to balance the trade-off between importance and diversity when ranking the visualizations (**Section 3**).
- We propose the novel *DiVE* schemes, that employs various algorithms to evaluate the recommended visualizations based on the hybrid ranking/objective function (**Section 4**).
- We present optimization techniques that leverage the hybrid objective function to substantially reduce the computational costs (**Section 4**).
- We conduct an extensive experimental evaluation on real datasets, which compare the performance of various algorithms and illustrate the benefits achieved by *DiVE* both in terms of effectiveness and efficiency (**Section 5**).

2 PRELIMINARIES AND RELATED WORK

Several recent research efforts have been directed to the challenging task of recommending aggregate views that reveal interesting data-driven insights (e.g., [3, 15, 16]). As in previous work, we assume a similar model, in which a visual data exploration session starts with an analyst submitting a query Q on a multi-dimensional database D_B. Essentially, Q selects a subset D_Q from D_B by specifying a query predicate T. Hence, Q is simply defined as: Q: `SELECT *` `FROM` D_B `WHERE` T;

Ideally, the analysts would like to generate some aggregate views (e.g., bar charts or scatter plots) that unearth some valuable insights from the selected data subset D_Q. However, achieving that goal is only possible if the analyst knows exactly what to look for! That is, if they know the parameters, which specify some aggregate views that lead to those valuable insights (e.g., aggregate functions, grouping attributes, etc.). Meanwhile, such parameters only become clear in "hindsight" after spending long time exploring the underlying database. Hence, the goal of existing work, such as [3, 8, 15–17], is to *automatically* recommend such aggregate views.

To specify and recommend such views, as in previous work, we consider a multi-dimensional database D_B, which consists of a set of dimensional attributes $\mathbb{A}$ and a set of measure attributes $\mathbb{M}$. Also, let $\mathbb{F}$ be a set of possible aggregate functions over measure attributes, such as `COUNT`, `AVG`, `SUM`, `MIN` and `MAX`. Hence, specifying different combinations of dimension and measure attributes along with various aggregate functions, generates a set of possible views $\mathbb{V}$ over the selected dataset D_Q. For instance, a possible aggregate view V_i is specified by a tuple $<A_i, M_i, F_i>$, where $A_i \in \mathbb{A}$, $M_i \in \mathbb{M}$, and $F_i \in \mathbb{F}$, and it can be formally defined as: V_i: `SELECT` A_i, F_i (M_i) `FROM` D_B `WHERE` T `GROUP BY` A_i;

Clearly, an analyst would be interested in those views that reveal interesting insights. However, manually looking for insights in

each view $V_i \in \mathbb{V}$ is a labor-intensive and time-consuming process. For instance, consider again our example in the previous section. In that example, let D_B be the Cleveland Heart Disease table (i.e., `tb_heart_disease`) and the analyst is selecting the subset of patients with heart disease (i.e., D_Q = `disease` subset). Hence, the number of views to explore is equal to: $|\mathbb{V}| = |\mathbb{A}| \times |\mathbb{M}| \times |\mathbb{F}|$, where $|\mathbb{F}|$ is the number of SQL aggregate functions, and $|\mathbb{A}|$ and $|\mathbb{M}|$ are the number of attribute and measures in `tb_heart_disease`, respectively. For that medium-dimensionality dataset, that value of $|\mathbb{V}|$ goes up to 180 views, which is clearly unfeasible for manual exploration. Such challenge motivated multiple research efforts that focused on automatic recommendation of views based on some metrics that capture the utility of a recommended view (e.g., [3, 5, 8, 10–12, 15–17]. Some of those works focus on recommending visualizations to facilitate a particular user intent or task. For example: using user feedback as a basis for view recommendation [8], recommending interactive visualizations on the webpage and engage user for collaboration and discussion [17], explanations for a certain behavior, finding data anomalies or outliers and correlations among data attributes [5, 12]. Hence, the criteria for ranking the visualizations is driven by the user intent. However, in visual data exploration, often the intent of the user is not clear. Towards that end, data driven metrics are employed to capture the interestingness or importance of a recommended visualization.

Among the data driven metrics, recent case studies have shown that a *deviation-based* metric is effective in providing analysts with *important* visualizations that highlight some of the particular trends of the analyzed datasets [10, 11, 15, 16].

In particular, the deviation-based metric measures the distance between $V_i(D_Q)$ and $V_i(D_R)$. That is, it measures the deviation between the aggregate view V_i generated from the subset data D_Q vs. that generated from a reference dataset D_R, where $V_i(D_Q)$ is denoted as *target* view, whereas $V_i(D_R)$ is denoted as *reference* view. That reference dataset could be the whole database (i.e., $D_R = D_B$) or a selected subset of the database (e.g., Sec. 1). The premise underlying the deviation-based metric is that a view V_i that results in a high deviation is expected to reveal some important insights that are very particular to the subset D_Q and distinguish it from the patterns in D_R. In case, $D_R = D_B$, then the patterns extracted from D_Q are fundamentally different from the generals ones manifested in the entire database D_B.

While recommending views based on their importance has been shown to reveal some interesting insight, it also suffers from the drawback of recommending similar and redundant views, which leaves the data analyst with a limited scope of the possible insights. The example has been shown in Sec. 1, two views which generated by same attribute and meassure and different aggregate function, both of them are considered as important views, however, those are similar and delivering limited amount of gained insights. To address that limitation, in this work we posit that employing *diversification* techniques in the process of view recommendation allows eliminating that redundancy and provides a good and concise coverage of the possible insights to be discovered. In the next section, we discuss in details the formulation of both importance and diversity, and their impact on the view recommendation process.

Table 1: Table of Symbols

Symbol	Description
k	no. of top recommended views
S	set of top-k recommended views
$\mathbb{V}$	set of all possible views
X	set of all candidate views
A	a dimensional attribute
M	a measure attribute
F	aggregate function
Q	a user query
D_B	a multi-dimensional database
D_Q	a target subset of D_B
D_R	a reference subset of D_B
V_i	a view query
$I(V_i)$	importance score of V_i
$I(S)$	importance score of views in S
$f(S,D)$	diversity score of views in S
$F(S)$	hybrid objective utility function value of S
$U(V_i)$	the utility score of each candidate view

3 DIVERSIFYING RECOMMENDED VISUALIZATIONS

In order to diversifying recommended visualizations, we work at two levels. At the first level, content driven deviation metric is used to evaluate that how "important" is the content of the view as compared to the reference view. At the second level, we evaluate contextually how different a view is from other views in the recommended set using context driven deviation measure. The details of those two level are described next.

3.1 Content-Driven Deviation

As briefly described in the previous section, in this work we adopt a deviation-based metric to quantify the importance of an aggregate view [15, 16]. Essentially, the deviation-based metric compares an aggregate view generated from the selected subset dataset D_Q (i.e., target view $V_i(D_Q)$) to the same view if generated from a reference dataset D_R (i.e., reference view $V_i(D_R)$).

Clearly, the deviation between a target and a reference view is a *data-driven* metric. That is, it measures the deviation between the aggregate *result* of $V_i(D_Q)$ and that of $V_i(D_R)$. Consequently, and from a visualization point of view, that deviation is a *content-based* metric that captures the difference between the content of the visualization generated by $V_i(D_Q)$ vs. the visual content generated from $V_i(D_R)$. In the next, we formally describe the standard computation of that data-driven content-based metric, whereas the discussion of its counterpart context-driven metric is deferred to the next section.

To calculate the content-based deviation, each target view $V_i(D_Q)$ is normalized into a *probability distribution* $P[V_i(D_Q)]$ and similarly, each reference view into $P[V_i(D_R)]$. In particular, consider an aggregate view $V_i = <A_i, M_i, F_i>$. The result of that view can be represented as the set of tuples: $<(a_j, g_j), (a_j, g_j), ..., (a_t, g_t)>$, where t is the number of distinct values (i.e., groups) in attribute A_i, a_j is the j-th group in attribute A_i, and g_j is the aggregated

Figure 3: Content vs. Context of views.

value $F_i(M_i)$ for the group a_j [3, 15]. Hence, V is normalized by the sum of aggregate values $G = \sum_{j=1}^{t} g_j$, resulting in the probability distribution $P[V_i] = <\frac{g_1}{G}, \frac{g_2}{G}, ..., \frac{g_t}{G}>$.

Finally, the importance score of V_i is measured in terms of the distance between $P[V_i(D_Q)]$ and $P[V_i(D_R)]$, and is simply defined as:

$$I(V_i) = dist \left(\mathcal{P} \left[V_i(D_Q) \right], \mathcal{P} \left[V_i(D_R) \right] \right) \tag{1}$$

where $I(V_i)$ is the importance score of V_i and $dist$ is a distance function. Similar to existing work (e.g., [15, 16]), we adopt a Euclidian distance function, but other distance measures are also applicable (e.g., Earth Mover's distance, K-L divergence, etc.).

In current approaches for view recommendation, the importance value $I(V_i)$ of each possible view V_i is computed, and the k views with the highest deviation are recommended (i.e., *top-k*) (e.g., [15]). However, in this work, our goal is to ensure that recommended views provide a good coverage of possible insights, which is achieved by considering the context of the recommended views, which is described next.

3.2 Context-Driven Deviation

As mentioned above, recommending top-k views based only on their data content (i.e., content-driven deviation) often leads to a set of similar views. In order to provide full coverage of all possible interesting insights, in this work, we posit that achieving *diversity* within the set of recommended views is an essential quality measure. Diversity has been well known and widely used in recommendation systems for maximizing information gain and minimizing redundancy (e.g., [9, 18, 20, 21]). At a high level, diversity essentially measures how different (i.e., diverse) are the individual data objects within a set.

Before discussing the details of diversity computation in Sec. 3.3, it is important to notice that central to that computation is some

notion of distance measure between data objects. Existing work provides multiple metrics for measuring that distance between traditional data objects, such as web documents (e.g., [1, 9, 21]), database tuples (e,g,, [14]), etc. However, our work in this paper is the first to consider diversity in the context of aggregate data visualizations. As such, a metric is needed to capture and quantify the (dis)similarity between the distinct features of different visualizations. Meanwhile, as each visualization is merely a data view generated by an SQL aggregate query, such metric naturally lends itself to considering the query underlying each view. That is, the query that has been executed to create the visualization. In turn, the distance between two visualizations is measured based on the distance between their underlying queries. Hence, in addition to the data-driven content-based deviation described above, here we also introduce a query-driven *context-based* deviation metric.

To measure the context-based deviation between two visualizations, we simply measure the distance between their underlying queries. Towards this, we extend on existing work in the area of query recommendation and refinement (e.g., [6, 7, 14]). In that work, the distance between two range queries q_1 and q_2 is mapped to that of measuring the edit distance needed to transform q_1 into q_2, where the set of allowed transformation are: add, delete, or modify a predicate. In the context of our work, however, views are generated from aggregate queries without range predicates. In particular, a view is fully defined in terms of a combination of attribute, measure and an aggregate function. Hence, in addition to the content of a view V_i which is described by its probability distribution (i.e., $P(V_i)$) as defined in Sec. 3.1), we also consider the context of the view $E(V_i)$, which is defined in terms of the query underlying V_i as: $E(V_i) = [A_i, M_i, F_i]$.

Such definition of view context leads to a special case of the existing work on query recommendation (e.g., [6, 7, 14]), in which the normalized distance between two queries is simply measured

using the *Jaccard* similarity measure. Hence, the Jaccard similarity between two aggregate views V_i and V_j is measured as:

$$J(V_i, V_j) = \frac{|E(V_i) \cap E(V_j)|}{|E(V_i) \cup E(V_j)|}$$

We note that the jaccard similarity assigns equal weights to each of the element in a set. Accordingly, when applied to aggregate views, then two views with the same attribute and different measure and aggregate function will have the same similarity score as any other pair of views with same measure but different attribute and aggregate function. However, an analyst may consider two views with the same attribute A_i more similar than two views with same measure attribute M_i. To allow the analyst to specify such preference, each contextual component of a view is associated with a weight that specifies its impact on determining the (dis)similarity between views. Specifically, let w_i be the weight assigned to i^{th} element of set $E(V_i)$, where $\sum_{i=1}^{3} w_i = 1$. Then, the similarity between views V_i and V_j is measured as: $J(V_i, V_j) = \dfrac{\sum_{i \in V_i \cap V_j} w_i}{\sum_{i \in V_i \cup V_j} w_i}$

Consequently, the context-based deviation between V_i and V_j is calculated as:

$$D\left(V_i, V_j\right) = 1 - J\left(V_i, V_j\right) \tag{2}$$

As a summarize, the difference betweeen content and context of view is described in Figure 3. Content is the probability distribution of the aggregated query result, whereas, Context is described as a set containing the name of the attribute, measure and function used to generate the view.

3.3 Problem Definition

In this section, we formally define our problem for recommending diversified interesting aggregate views. Towards this, we first define the metrics to measure the performance of our proposed visualization recommendation system in terms of: 1) the quality of recommended visualizations, and 2) the processing cost incurred in computing those visualizations.

3.3.1 Hybrid Objective Function. Our hybrid objective function is designed to consider both the importance and diversity of the recommended views. Particularly, it integrates two components: 1) the total importance score of set S and 2) the diversity score of S.

The importance score of the a S is calculated as the average value of the importance measure of each view in S, as given in Eq.1. Hence, the total importance score of S is defined as:

$$I(S) = \sum_{i=1}^{k} \frac{I(V_i)}{I_u}, V_i \in S$$

where I_u is the upper bound on the importance score for an individual view, which is achieved when for each group a_i, the corresponding value $\frac{g_i}{G}$ in $P[V_i(D_R)]$ or $P[V_i(D_Q)]$ is zero. Thus, $I_u = \sqrt{2}$, and is used to normalize the average importance score for set S.

In order to measure the diversity of a set of objects, several diversity functions have been employed in the literature [1, 18]. Among those, previous research has mostly focused on measuring diversity based on either the average or the minimum of the pairwise distances between the elements of a set [19]. In this work, we focus on the first of those variants (i.e., average), as it maximizes the coverage of S. Hence, given a distance metric $D\left(V_i, V_j\right)$, as given

in equation 2, the diversity of a set S can be simply measured as follows:

$$f(S, D) = \frac{1}{k(k-1)} \sum_{i=1}^{k} \sum_{j>i}^{k} D\left(V_i, V_j\right), V_i, V_j \in S$$

Since the maximum context-based deviation between any two views in equation 2 is 1.0, then dividing the sum of distances by $k(k-1)$ ensures that the diversity score of set S is normalized and bounded by 1.0.

Next, we define our proposed hybrid objective function that captures both the importance and diversity of the set of recommended views S. Specifically, for a set of views $S \subseteq V$, our hybrid objective function is formulated as the linear weighted combination of the importance score, $I(S)$ and diversity score $f(S, D)$, and is defined as:

$$F(S) = (1 - \lambda) \times I(S) + \lambda \times f(S, D) \tag{3}$$

where $0 \leq \lambda \geq 1$ is employed to control the preference given to each of the importance and diversity components. For instance, a higher value of λ results in a set of more diverse views, whereas a lower value of λ generates a set of the most important views, which is likely to exhibit redundancy in the recommended views.

Given the hybrid objective function, our goal is to find an optimum set of views S^* that maximizes the objective function $F(S)$, which is defined as follows:

DEFINITION 1. **Recommending diversified important views** : *Given a target subset D_Q and a reference subset D_R, the goal is to recommend a set $S \subseteq \mathbb{V}$, where $|S| = k$, and $\mathbb{V}$ is the set of all possible target views, such that the overall hybrid objective $F(S)$ is maximized.*

Given the definition above, the quality of the recommended set of views is measured in terms of the value of the hybrid objective function $F(S)$.

$$S^* = \underset{\substack{S \subseteq \mathbb{V} \\ |S|=k}}{\operatorname{argmax}} F(S) \tag{4}$$

3.3.2 Cost of Visualization Recommendation. Exisiting research has shown that recommending aggregate data visualizations based on data-driven content-based deviation is a computationally expensive task [3, 15, 16]. Moreover, integrating diversification to the view recommendation problem, as described above, further increases that computational cost. In particular, the incurred processing cost includes the following two components:

(1) Query processing cost C_Q: measured in terms of the time needed to execute and compare all the queries underlying the set of target views as well as their corresponding reference views (i.e., content-based deviation).
(2) View diversification cost C_D: measured in terms of the time needed to compute all the pairwise distances between each pair of target views (i.e., context-based deviation).

Consequently, the total cost C_T for recommending a set of views is simple defined as:

$$C_T = C_Q + C_D$$

In principle, traditional data diversification methods that consider both relevance and diversity can be directly applied in the context of our problem to maximize the overall utility function formulated in Eq.3. For instance, in the context of recommending

web search, such methods are designed to recommend a set of diversified objects (e.g., web documents) that are relevant to the user needs. However, in that setting, the relevance of an object is either given or simply computed. To the contrary, in our setting for view recommendation, the importance of a view is a computational expensive operation, which requires the execution of a target and reference view. As such, directly applying those methods leads to a "process-first-diversify-next" approach, in which all possible data visualization are generated first via executing a large number of aggregate queries. To address that challenge and minimize the incurred query processing cost, in the next section we propose an integrated scheme called *DiVE*, leverages the properties of both the importance and diversity to prune a large number of a large number of low-utility views without compromising the quality of recommendations, as described next.

4 THE DIVE SCHEMES

In this section, we present our DiVE schemes which capture both the importance, as well as the diversity of the insights revealed by the recommended views. Before presenting the details of how DiVE evaluates the hybrid objective function, we first present two baseline for evaluating top-k views from exiting work.

4.1 Baseline Solutions

As baseline solutions to compare the performance of our proposed DiVE schemes, we simply incorporate methods from existing work that optimize either for importance or diversity. In terms of diversity, we employ the classical *Greedy Construction* algorithm [13], which has been shown to maximize diversity within reasonable bounds compared to the optimal solution [18, 20]. In this work, we refer to that baseline as *Greedy-Diversity*. Similarly, in terms of importance, we adopt the work on SeeDB for recommending the top-k views with the highest deviation [15, 16]. Particularly, in that method, all possible target and reference views are generated by executing their underlying queries, then the list of views is linearly scanned to recommend the top-k for which the target view shows high deviation from its corresponding reference view (denoted as *Linear-Importance* in this work). Clearly, those two methods are "oblivious" to our hybrid objective function (i.e., Eq.3). In particular, as shown in our experimental evaluation (Sec. 6), each of those two methods performs well under extreme settings of our hybrid function. As expected, Greedy-Diversity provides its best performance when $\lambda = 1.0$ (i.e., all preference is given to diversity), whereas Linear-Importance is the winner when $\lambda = 0.0$ (i.e., all preference is given to importance). In the following, we present our DiVE schemes which are able to provide the best performance (i.e., maximize the overall hybrid objective), irrespective of the value of λ, while minimizing the processing time.

4.2 The DiVE-Greedy Scheme

In this section, we discuss our first DiVE scheme (*DiVE-Greedy*), which simply extends the basic Greedy Construction algorithm to work under our hybrid objective function (i.e., Eq. 3). Such extension is straightforward and is described in Algorithm 1. Similar to the classical Greedy Construction, DiVE-Greedy initializes the set S with the two most distant views, where the distance between any

---

Algorithm 1: *DiVE* Greedy

Input: Set of views V and result set size k
Output: Result set $S \geq V$, |S| = k

1 $S \leftarrow \left[V_i, V_j\right]$ get two most distant views;
2 $X \leftarrow [V \backslash S]$;
3 $i \leftarrow len\,(S)$;
4 **while** $|S| < k$ **do**
5 $\quad$ $v_i \leftarrow argmax\,(1 - \lambda) \times I\,(v_i) + \lambda \times setDist\,(v_i, S)$;
6 $\quad$ $S.add\,(v_i)$;
7 $\quad$ $X.remove\,(v_i)$;
8 **end**
9 return S

---

two views is calculated using our context-based function, as given in Eq.2. Then, DiVE-Greedy iteratively selects new views to be added to S. Particularly, in each iteration a view is selected from the set of remaining views X and is added to S. To make that selection, DiVE-Greedy assigns a score to each view in X, which is based on the hybrid objective function $F\,(S)$, as defined in equation 3. Specifically, the utility score assigned to a view $V_i \in X$ is computed as:

$$U\,(V_i) = (1 - \lambda) \times I\,(V_i) + \lambda \times setDist\,(V_i, S) \qquad (5)$$

where

$$setDist\,(V_i, S) = \frac{1}{|S|} \sum_{\substack{j=1 \\ V_j \in S}}^{|S|} D\left(V_i, V_j\right)$$

Thus, in each iteration, the view with highest utility score is selected and added to S, until $|S| = k$, as shown in Algorithm 1. ,

DiVE-Greedy Cost: We note that the only difference between our DiVE-Greedy scheme and our baseline Greedy-Diversity (i.e., the classical Greedy Construction algorithm) is in the utility score assigned to each view (i.e., $U(V_i)$ in Eq.5). In fact, in the special case where $\lambda = 1.0$, Eq. 5 boils down to $U\,(V_i) = setDist\,(V_i, S)$, which is the same score used by Greedy-Diversity for maximizing diversification. However, that simple change in the utility score leads to executing the query underlying each view V_i in order to compute the $(1 - \lambda) \times I\,(V_i)$ component of its score. Hence, the overall cost of DiVE-Greedy is $C_T = C_Q + C_D$, as opposed to the cost of Greedy-Diversity, which is only $C_T = C_D$, where C_Q is the query processing cost (i.e., data-driven), and C_D is the cost for computing Jaccard distances (i.e., query-driven), as described in Sec. 3.

Clearly, C_Q is equal to the number of possible views and is $O(n)$, where n is the number of possible views, whereas C_D is $O(kn)$, where k is the number of recommended views. Hence, as presented, our extended DiVE-Greedy is an instance of what has been described earlier as a "process-first-diversity-next" approach [], I am not sure which papers should I cited here in which diversification only takes place after all possible views are executed to calculate their importance. In the next section, we describe our work to overcome the limitations of that approach.

4.3 DiVE-Greedy-Static Pruning

As mentioned in the previous section, the cost of the DiVE-Greedy algorithm is dominated by the query processing cost C_Q and is proportional to the number of possible views. Although hundreds of views are generated for a given subset of data D_Q, only a small fraction of those views are actually of interest and are candidates to be included in the top-k set. As such, a significant fraction of the query processing cost is incurred in evaluating low-utility views. This observation motivated us to propose a pruning technique to minimize the search space of views and narrow it down to the most promising ones, as follows.

Our proposed pruning technique is based on the observation that the utility score of each view $U(V_i)$ is a weighted sum of two measures; 1) the importance score of the view (i.e., $I(V_i)$), and 2) the distance of the view from S (i.e., $setDist(V_i, S)$). We note that the computation of $setDist(V_i, S)$ is a CPU-bound requires fast operation. To the contrary, computing the importance score of a view $I(V_i)$ is an expensive operation that requires executing two queries to generate the target and reference data for V_i. Thus, we employ a *max-min* pruning technique, that leverages the diversity score to bound the maximum utility score achieved by each view V_i, and in turn allows for pruning low-utility views without incurring the high cost for evaluating their importance.

Particularly, max-min utilizes the maximum and minimum bound on the importance score $I(V_i)$. Clearly, the minimum utility score is $I_0 = 0$, whereas the maximum utility value I_u a view can score is computed as $I_u = \sqrt{2}$ (Sec. 3.3). Using those bounds, in each iteration the maximum $maxU(V_i)$ and minimum utility $minU(V_i)$ score for each view $V_i \in X$ is calculated as:

$$maxU(V_i) = (1 - \lambda) \times I_u(V_i) + \lambda \times setDist(V_i, S)$$
$$minU(V_i) = (1 - \lambda) \times I_0(V_i) + \lambda \times setDist(V_i, S)$$

Accordingly, the maximum value of $minU$ is recorded (i.e., maximum of mimumum). Then, if $maxU$ of any candidate view is less than the maximum of $minU$, then that view is pruned.

4.4 DiVE-Swap Scheme

Naturally, our DiVE-Greedy scheme described above will achieve a high pruning power for high values of λ (i.e., more emphasize is given to the diversity score $setDist(V_i, S)$). However, that pruning power diminishes for low values of λ (i.e., more emphasize is given to the importance score $I(V_i)$). That discrepancy in performance happens because of two reasons: (1) the constructive and iterative nature of the greedy algorithm, and (2) utilizing the upper bound on importance I_u when computing $maxU(V_i)$. To explain those reasons, recall that the amount of pruning is shaped by the maximum value of $minU(V_i)$. Now, consider an extreme case, in which $\lambda = 0.0$. In that case, all values of $minU(V_i)$ will also be equal to zero, and no pruning will take place. Similarly, for low values of $\lambda > 0$, the maximum value of $minU(V_i)$ will still be low, leading to limited pruning. That pruning is further reduced when combined with a high value of I_u as it leads to all views acquiring high $maxU(V_i)$ and cannot be pruned. To address those limitations, we propose an adaptive scheme for setting I_u, which is described in Sec. 4.6, and we orthogonally introduce our *SWAP*-based DiVE scheme, which is described next.

Algorithm 2: *DiVE* Swap

Input: Set of views V and result set size k
Output: Result set $S \geq V$, |S| = k

1 $S \leftarrow$ Result set of only importance or only diversity;
2 $X \leftarrow [V \setminus S]$;
3 $F_{current} \leftarrow 0$;
4 $improve \leftarrow True$;
5 **while** *improve = True* **do**
6   **for** *i in set X* **do**
7          $S' \leftarrow S$;
8   **for** *j in set S* **do**
9   **if** $F(S') < F(S \setminus S[j] \cup X[i])$ **then**
10                 $S' \leftarrow S \setminus j \cup X[i]$;
11   **end**
12   **end**
13   **if** $F(S') > F(S)$ **then**
14             $S \leftarrow S'$
15   **end**
16   **end**
17   **if** $F(S) > F_{current}$ **then**
18         $F_{current} \leftarrow F(S)$;
19         $improve \leftarrow True$;
20   **else**
21         $improve \leftarrow False$;
22   **end**
23 **end**
24 return S

The DiVE-Greedy algorithm presented in the previous section is of the constructive type. That is, it starts without a initial set of views and incrementally constructs it by adding one view at a time. To the contrary, our DiVE-SWAP presented in this section falls under the *local search* type of algorithms. In general, a local search algorithm starts out with a complete initial solution and then attempts to find a better solution in the neighborhood of that initial one. Like constructive algorithms, local search algorithms are also widely used in solving optimization problems including diversification. For instance, the SWAP local search method has been utilized to maximize diversity [2, 18], and in this paper, we further expand into our DiVE scheme recommending aggregate views.

The basic idea underlying DiVE-SWAP is to start with an initial set S of size k and then iteratively modify the set S in order to im-prove the value of the objective function $F(S)$. One of the main design criteria in local search algorithms is the choice of the initial solution. In DiVE-SWAP, we consider two variants: 1) *DiVE-iSwap*, and 2) *DiVE-dSwap*. In DiVE-iSwap, S is initialized with the k views that maximize importance (i.e., all preference is given to importance) and can be easily obtained using our baseline Linear-Importance (Sec. 4.1). Alternatively, in DiVE-dSWAP, will we use dSWAP instead of dSwap?S is initialized with the k views that maximize diversity (i.e., all preference is given to diversity), which also can be easily obtained using our baseline Greedy-Diversity (Sec. 4.1).

Apart from the initialization approach, both variants work similarly. Particularly, in each iteration, each unselected view $X_i \in X$ is interchanged with all the selected views in $S_j \in S$ (Algorithm 2 line 9). That is, the overall hybrid objective is computed for $F(S \backslash S_j \cup X_i)$. Then the one interchange that leads to the highest new value for F is applied and S is updated accordingly (Algorithm 2 line 10). Such iterations are repeated until no more views can be swapped between X and S, which is reached when no further improvement is achieved in the value of F.

In comparison to DiVE-Greedy, DiVE-SWAP incurs the same query processing cost C_Q. Furthermore, its incurs even higher C_D cost for computing diversity, which can reach up to $O\left(kn^2\right)$ in the worst case I have checked this in Zaeem paper. However, as described in the next section, DiVE-SWAP offers a valuable opportunity for maximizing the number of pruned views, and in turn reducing the query processing cost C_Q, which is the dominant factor in determining the efficiency.

4.5 DiVE-dSwap Static Pruning

While both DiVE-iSwap and DiVE-dSwap described above incur the same cost, they offer substantially different performance when combined with pruning techniques. Particularly, consider DiVE-iSwap, which is initialized with the k views that maximize importance. To select those views, all possible views have to be generated first, which in turn requires processing all their corresponding queries. Hence, DiVE-iSwap simply eliminates all opportunities for pruning as all views are executed in the initialization phase. To the contrary, DiVE-dSwap is initialized with the k views that maximize diversity. To select that initial set, no query execution is needed and the processing is limited to computing the context-based deviation distances, which is significantly low compared to query processing. Hence, DiVE-dSwap allows integrating effective pruning techniques, as described next.

Recall that under DiVE-dSwap, in each iteration a view $X_i \in X$ is selected to replace a view $S_j \in S$. The criterion for that selected view is to improve the overall hybrid objective F. That is, $F(S \backslash S_j \cup X_i) > F(S)$. Hence, the task boils down to finding that pair of views $< X_i, S_j >$ that provide the maximum improvement in F once interchanged. Without pruning, that requires iterating through S and X simultaneously and computing F for each pair, which requires processing and generating each view in X. To avoid such expensive processing and enable pruning, the following steps are taken. A list is created for all possible swap pairs $< X_i, S_j >$. That list is sorted based on the diversity achieved if the swap is to be made. That is, eq. Hi Mohamed, do you mean new equation? If pseudocode I can wrote that, but equation I was trying but confused Notice that up to this point the only processing needed is to compute diversity without any query execution to evaluate the importance of any X_i. Given that setting, the task is similar to DiVE-Greedy-Static.

To find the top-1 view, each view X_i is initially assigned an importance equal to the upper bound I_u. In turn, the upper bound of F achieved by X_i is computed as: $maxF(S \backslash S_j \cup X_i)$, which is based on the actual diversity achieved by the swap, and the upper bound on importance. As such, $maxF(S \backslash S_j \cup X_i)$ is compared against $F(S)$, leading to one of the following two cases: If $maxF(S \backslash S_j \cup X_i) > F(S)$, then the swap $< X_i, S_j >$ can "potentially" improve F. Hence, at that stage the view X_i needs to be generated in order to evaluate its actual importance $I(X_i)$. Otherwise, the pair $< X_i, S_j >$ is pruned if:

$$maxF(S \backslash S_j \cup X_i) < F(S)$$

Simply put, if the upper bound $maxF$ achieved by that swap is still less than the current $F(S)$, then the actual $F(S \backslash S_j \cup X_i)$ is guaranteed to be less than $F(S)$ and the pair $< X_i, S_j >$ can be safely ignored (i.e, pruned). More importantly, since the list of pairs is sorted by diversity, then the next views are also guaranteed to provide no improvement and that iteration of DiVE-SWAP reaches *early termination*. Hence, for all the remaining views no query processing is needed, which significantly reduces the overall cost.

At this point, it is especially important to examine and contrast the pruning power achieved by each of DiVE-Greedy and DiVE-SWAP. For DiVE-Greedy, recall that pruning is attained for those views where: $maxU(V_i) < maxminU$. However, since DiVE-Greedy is a constructive algorithm, the set S is incrementally constructed iteration by iteration until $|S| = k$. Hence, in the first iterations S has a very small number of selected views with a minimum of $|S| = 2$. Naturally, when S is a small set, then most of the remaining unselected views in X are expected to exhibit high diversity, since the majority of them will be very dissimilar from the small set of views in S. Accordingly, for most of the views in X, the value $maxU(V_i)$ will be relatively high, due to achieving high score on diversity. Hence, most views will fail to satisfy the pruning condition and are consequently executed incurring high query processing cost.

To the contrary, DiVE-SWAP is initiated with a set S of k diverse views. Hence, it will initially have a reasonably high $F(S)$. Moreover, many views in X will be more or less "close" to some view in S. Hence, the swaps that involve those views will score low on diversity, and in turn low $maxF$. Since pruning happens when $maxF(S \backslash S_j \cup X_i) < F(S)$, the combination of those two factors above (i.e., high F and low $maxF$) allows for many views satisfying the pruning condition, which improves the pruning power of DiVE-SWAP over DiVE-Greedy. That pruning power can be further improved by relaxing the assumption about maximum importance I_u, as described next.

4.6 Adaptive Pruning Scheme

In general, both the pruning schemes provided by DiVE-Greedy and DiVE-SWAP rely on the fundamental idea of evaluating the upper bound of the benefit provided by a view V_i towards the objective F. If that maximum benefit is still not enough to consider V_i to join S, then V_i is pruned and its query processing cost is saved. Moreover, to evaluate that upper bound, both schemes compute the actual diversity offered by V_i and instead of computing its actual importance, it is substituted with the maximum attainable important I_u. Naturally, overestimating $I(V_i)$ leads to overestimating its benefit and consequently limited pruning power is achieved. Meanwhile, for most datasets I_u is in fact an overestimation of $I(V_i)$. Hence, our goal in this section to provide a tighter bound on $I(V_i)$, which allows for maximum pruning while maintaining the quality of the solution.

Recall that I_u is achieved when for each group a_i, the corresponding value $\frac{g_i}{G}$ in $P[V_i(D_R)]$ or $P[V_i(D_Q)]$ is zero. Hence, I_u is

a theoretical bound for the maximum importance achieved by any view in any dataset. For most real datasets, however, that condition is rarely satisfied and the actual upper bound I_{au} is typically much smaller than I_u. Meanwhile, a hypothetical pruning scheme that utilizes that actual upper bound I_{au} is expected to deliver more pruning power than the schemes using the theoretical upper bound I_u, especially when $I_{au} \ll I_u$. In practice, however, that hypothetical scheme is not achievable since obtaining the value I_{au} requires executing all the possible views, which is clearly in conflict with the goal of pruning.

Accordingly, rather than using overestimated I_u or obtaining the actual I_{au}, our goal is to estimate I_{au} with high accuracy and minimum number of query executions. In particular, given the set of possible views $\mathbb{V}$, the goal is to estimate the maximum importance $\bar{I}_{au}$ given by some view in $\mathbb{V}$. However, estimating the maximum value of a population is known to be a challenging problem, as opposed to estimating other statistics such as average or sum [4]. That challenge is even emphasized when the values exhibited by the population are skewed and do not follow a typical normal distribution, which is typically the case for the importance value of views. Thus, instead of estimating I_{au}, we rely on *non-parametric predictive interval* models to determine its value with certain level of confidence without any assumption on the population [4]. To apply that model, some sample views are executed and the maximum importance observed in that sample is recorded as $\bar{I}_{au}$. To determine the number of samples, a *Predictive Interval (PI)* is to be defined, such that: $PI = \dfrac{(m-1)}{(m+1)}$, where m is the number of samples. For instance, setting $m = 19$, results in prediction interval PI = 90%. That is, 90% of the time, the importance value of an unseen view V_i will be less than the maximum importance seen so far. Clearly, the higher the PI value, the higher the accuracy of $\bar{I}_{au}$ in estimating I_{au}, but also requires executing more views. In this work, we find that a value of $PI = 95\%$ in the experiment, we used 97 and 80, but here you said 95, should we add plot for 95? is able to balance the tradeoff between minimizing the number of executed queries and maximizing the objective value F, as illustrated in the next sections.

5   EXPERIMENTAL TESTBED

In this section, we present an evaluation of our proposed *DiVE* scheme both in terms of effectiveness and efficiency when returning diversified interesting visualizations.Table 2 summarizes the different parameters used in our experimental evaluation. All the experiments . The default values are presented in bold. All experiments were run on a single machine with 16 GB RAM and a 64 bit, Intel Core i7-7700 processor. All the performance measures are averaged over ten runs. We have conducted our experiments over following real datasets:

- Heart Disease Dataset [2]: This dataset is comprised of 9 dimensional attributes and 5 measure attributes. The number of possible views are 180 with four aggregate functions.
- Airline (Flights) Dataset [3]: This dataset is comprised of 7 dimensional attributes and 4 measure attributes. The number of possible views are 112.

---

[2]http://archive.ics.uci.edu/ml/datasets/heart+Disease
[3]http://stat-computing.org/dataexpo/2009/the-data.html

Table 2: Parameters testbed in the experiments

Parameter	Range (**default**)
datasets	**Heart disease**, Flights
sample queries	**10**
diversity weight ratio	$3(A) : 2(M) : 1(F)$
tradeoff weight λ	0.0, 0.2, 0.4, **0.5**, 0.6, 0.8, 1.0
result set (size of k)	**5**, 15, 25, 35
prediction interval %	80 , 85, 90, 95, **97**, 98

In particular, we evaluate the performance of the following schemes:

- Linear-Importance: Selects top-k views on the basis of only the importance score.
- Greedy-Diversity: Selects top-k diverse views.
- DiVE-Greedy: Selects top-k views on the basis of hybrid objective function using greedy algorithm.
- DiVE-iSwap: Selects top-k views on the basis of hybrid objective function using swap algorithm initialized by Linear-importance method.
- DiVE-dSwap: Selects top-k views on the basis of hybrid objective function using swap algorithm initialized by Greedy-Diversity method.
- Static-pruning: DiVE-Greedy and DiVE-dSwap methods using static pruning technique to reduce the number of view queries as presented in section **??**
- Adaptive-pruning: DiVE-Greedy and DiVE-dSwap methods using adaptive pruning method as discussed in section **??**.

For each experiment a query workload of ten random queries is submitted to select ten different subsets of data D_Q. The performance measures are averaged over views recommended for each of those ten subsets. In particular, the performance of each scheme listed above is measured based on the following metrics:

- **Interestingness of the views:** measured as the value of hybrid objective function $F(S)$ for selected set of views S, as defined in equation 3.
- **Total Cost:** measured as the sum of the cost to execute view queries and the cost to diversify the views that is : $C_T = C_Q + C_D$

6   EXPERIMENTAL EVALUATION

In this section, we evaluate the sensitivity of *DiVE* scheme to the different parameters as given in Table 2.

6.0.1   **The impact of λ on the objective function $F(S)$.** The value of λ balances the trade off between importance and diversity score. Figure 4 shows how the performance of each scheme in terms of $F(S)$ is effected as the value of λ varies from 0 to 1. It is clearly seen in Figure 4, that for the lower values of λ the highest objective function value is achieved by Linear-Importance method. To the contrary, the Greedy-Diversity method achieves highest values of $F(S)$ as the λ approaches 1. Hence, there is a crossover between Linear-Importance and Greedy-Diversity. However, our proposed schemes based on the hybrid objective function have stable performance for all values of λ and are able to achieve $F(S)$ values higher

(a) Heart disease dataset

(b) Flights dataset
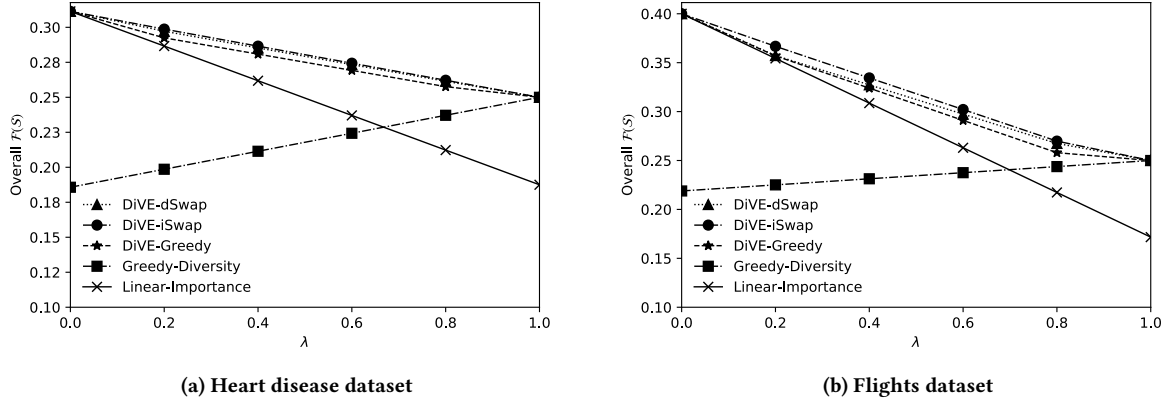
Figure 4: Impact of λ to overall objective function value $F(S)$ while k = 5 and running on real datasets

(a) Heart disease dataset

(b) Flights datasets

Figure 5: Overall objective function value $F(S)$ using different value of k and running on real datasets

than those achieved by Linear-Importance and Greedy-Diversity. compare the DiVE schemes, why value of F goes down for all with increase in λ

6.0.2 **The impact of k on the objective function** $F(S)$**.** Figure 5 shows the $F(S)$ values for various schemes as the number of recommended views k varies from 5 to 35. Overall the $F(S)$ value decreases with increasing value of k for all the schemes. This is because both average importance score and diversity of a set S decreases as new views are added to S. The views added earlier to S have higher importance score then the views added later. Similarly, the diversity function exhibits a diminishing marginal gain trend as the size of set S increases. The important observation here is the fact that our *DiVE* schemes always have higher overall objective function values as compared to the two extreme baselines approaches for all values of k. Among the *DiVE* schemes, DiVE-iSwap and DiVE-dSwap perform better than the Greedy versions. This is because, swap algorithm performs number of additional iterations to improve the value of the objective function.

Figure 6: Total time (seconds) to execute schemes on flights dataset using k =5 and $\lambda = 0.5$. It shows that costs are dominated by query costs C_Q
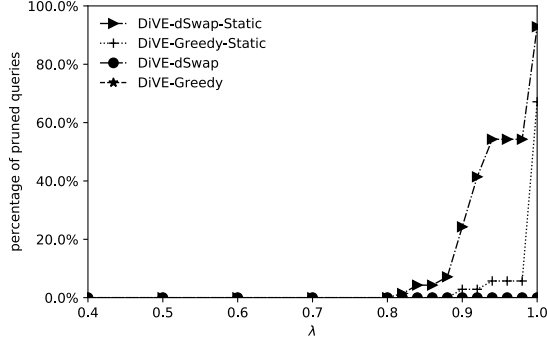
Figure 7: Performance of static pruning compared to without pruning in terms of cost, DiVE-dSwap-Static prunes queries around 20% while $\lambda = 0.9$


Figure 9: Impact of λ to the pruned queries of *DiVE-dSwap-Adaptive* scheme, running on Heart disease dataset using k = 5. Adaptive pruning schame able to prune queries start from $\lambda = 0.1$, it has more pruned queries while λ increased

Figure 8: Impact of λ to the pruned queries of *DiVE-Greedy-Adaptive* scheme in different used PI, running on Heart disease dataset using k = 5.

Figure 10: Impact of λ to the performance of *DiVE-dSwap-Adaptive* with different of PI in terms of the effectiveness, running on Heart disease dataset using k = 5. Using PI 80 can reduce the overall objective function $F(S)$ / the quality of the result

6.1 Execution time evaluation

In this experiment we have performed the cost analysis of *DiVE* schemes in terms of the execution time. Figure 6 plots the execution time for flights data set with $k = 5$ and $\lambda = 0.5$. The total execution time is split into the query execution time C_Q and the diversification cost C_D. It is clear from Figure 6 that the total execution time C_T is dominated by the cost of generating the views C_Q. Hence, the minimum cost is incurred by the Greedy-Diversity method which does not compute the view content to determine the importance score. For all other methods, the C_Q component of the execution time is same as all views are generated only once. However, the cost of diversification C_Q is slightly higher for DiVE-iSwap and DiVE-dSwap as compared to the DiVE-Greedy due the higher number of iterations.

6.2 Impact of static Pruning

In this experiment we present the performance of our proposed pruning techniques in terms of the number of pruned queries required to generate the views. The higher number of pruned queries result in the higher cost savings in the total query execution time.

Figure 7 shows the performance of our static pruning technique using the maximum value of importance score *maxI*. The evaluation is performed on the heart disease data set. Since, *maxI* value can be far from the actual importance scores of individual views, the percentage of pruned queries is 0 for lower values of λ. Only for λ close to 0.9 some queries get pruned. DiVE-dSwap-Static is able to prune higher numbe of queries than DiVE-Greedy-Static. At $\lambda = 1$ DiVE-dSwap-Static prunes 80% queries while DiVE-Greedy-Static prunes almost 65% queries. This is because for $\lambda = 1$ the weight of importance score is 0 in the hybrid objective function and the overall value of $F(S)$ is determined by the diversity score only.

6.3 Impact of Adaptive Pruning

In this experiment we analyze the performance of adaptive pruning technique. For clarity of presentation, we have shown the performance of DiVE-Greedy and DiVE-dSwap separately under different values of λ and prediction interval PI in Figure 8 and Figure 9 respectively.

DiVE-Greedy-Adaptive: As shown in Figure 8 , DiVE-Greedy-Adaptive is able to prune queries even for lower values of λ. The number of queries pruned increase significantly for higher values of λ. In comparison to the static pruning as shown in Figure 7, with adaptive pruning DiVE-Greedy is able to prune almost 20% more queries at $\lambda = 0.8$. The best performance is shown by DiVE-Greedy-Adaptive(0.80) as it executes a smaller number of sample queries.

DiVE-dSwap-Adaptive: Figure 9 shows the performance of DiVE-dSwap-Adaptive with different values of PI. In comparison to DiVE-Greedy-Adaptive, the number of pruned queries by DivE-dSwap-Adaptive are much higher for all values of λ. The interesting observation is the fact that DiVE-dSwap-Adative is able to prune 12% queries for λ values as low as 0.2. For higher values of λ the percentage of pruned queries is between 60% and 90%. Similar to DiVE-Greedy-Adaptive, highest number of queries are pruned for $PI = 0.8$. The performance declines slightly as we increase PI to 0.97. This is because when the sample size is large, many queries are already executed and the margin of cost savings by pruning the remaining queries is small.

Further we evaluate the effectiveness of methods with adaptive pruning in terms of the $F(S)$ values. Although, adaptive pruning is based on approximation of the actual importance score bounds, still the loss in the $F(S)$ values of the computed set of views is very small. Figure 10 shows the loss in $F(S)$ in comparison to the $F(S)$ values achieved by Hypothetical methods. The loss for both DiVE-Greedy-Adaptive and DiVE-dSwap-Adaptive is 0% for $PI = 0.97$, With a larger sample size the accuracy of approximated importance score is higher. However, for a smaller sample size of $PI = 0.80$ there is a maximum loss of 10% at $\lambda = 0.1$. The loss in $F(S)$ values decrease as λ increases as the impact of importance score becomes smaller in the hybrid objective function.

6.4 The impact of query subset on pruning

. The pruning power of our algorithms is effected by the maximum importance score views generated over a data subset D_Q in comparison to the reference subset D_R. If the highest importance score of any view generated over D_Q is far less than the upper bound of $\sqrt{2}$, then only few queries will be pruned. Similarly, for adaptive pruning technique higher variation in the importance scores of views will effect the accuracy of the maximum importance score estimated from the sample. Thus, for this experiment, we created three different subsets of heart disease dataset. And categorized them as low, medium and High. Where low dataset is the subset with low values of importance score and high is the subset with higher values of importance score. We executed random queries over those subsets and compared the performance of pruning methods in terms of number of queries pruned. The PI value is set to default value of 0.97. It can be seen in Figure 11 that both DiVE-Greedy-Adaptive and DiVE-dSwap-Adaptive perform better on low importance query load. For high importance query load less number of queries are pruned. This shows that the characteristic of the query subset D_Q has an impact on the efficiency of the pruning methods.

7 CONCLUSIONS

Figure 11: Number of pruned queries while using low vs. high importance query load in different value of λ, k = 5 and running on Heart disease dataset

In this paper we proposed an effective visualization recommendation scheme that promotes both importance and diversity in the recommended views. Our proposed scheme called *DiVE*, integrates importance and diversity into a *hybrid utility function* to provide good coverage of the possible insights to be discovered. In addition to employing a hybrid utility function for effective view recommendation, *DiVE* also leverages the properties of both the importance and diversity metrics to prune a large number of query executions without compromising the quality of recommendations. Hence, providing an efficient solution for a computationally expensive view recommendation task. We conducted extensive experimental evaluation on real data sets, which illustrate the benefits achieved by *DiVE* both in terms of effectiveness and efficiency.

REFERENCES

[1] Charles L.A Clarke and et al. 2008. Novelty and diversity in information retrieval evaluation. *SIGIR*.
[2] M Drosou and et al. 2010. Search Result Diversification. *SIGMOD Record* 39, 1.
[3] H Ehsan, M. Sharaf, and Panos K. Chrysanthis. 2016. MuVE: Efficient Multi-Objective View Recommendation for Visual Data Exploration. *ICDE*.
[4] Ying Hu and et al. 2009. Estimating Aggregates in Time-constrained Approximate Queries in Oracle. In *EDBT*. ACM, 1104–1107.
[5] Sean Kandel and et al. 2012. Profiler: Integrated statistical analysis and visualization for data quality assessment. In *AFI*. ACM, 547–554.
[6] V. Kantere. 2016. Query Similarity for Approximate Query Answering. In *DEXA*.
[7] Verena Kantere, George Orfanoudakis, Anastasios Kementsietsidis, and Timos K. Sellis. 2015. Query Relaxation across Heterogeneous Data Sources. In *CIKM*.
[8] Alicia Key, Bill Howe, Daniel Perry, and Cecilia R. Aragon. 2012. VizDeck: self-organizing dashboards for visual analytics. *SIGMOD Conference*.
[9] Davood Rafiei, Krishna Bharat, and Anand Shukla. 2010. Diversifying web search results. *WWW*.
[10] Thibault Sellam and Martin L. Kersten. 2016. Fast, Explainable View Detection to Characterize Exploration Queries. In *SSDBM*. 20:1–20:12.
[11] Thibault Sellam and Martin L. Kersten. 2016. Ziggy: Characterizing Query Results for Data Explorers. *PVLDB* 9, 13 (2016), 1473–1476.
[12] Jinwook Seo and Ben Shneiderman. 2006. Knowledge Discovery in High-Dimensional Data: Case Studies and a User Survey for the Rank-by-Feature Framework. *TVGC* 12, 3 (2006), 311–322.
[13] Barry Smyth and Paul Mcclave. 2001. Similarity vs . Diversity.
[14] Quoc Trung Tran and Chee-Yong Chan. 2010. How to ConQueR why-not questions. In *SIGMOD*.
[15] M Vartak and et al. 2015. SEEDB : Efficient Data-Driven Visualization Recommendations to Support Visual Analytics. *VLDB* 8.
[16] Manasi Vartak and Samuel Madden. 2014. S EE DB : Automatically Generating Query Visualizations. *VLDB*.
[17] Fernanda B. Viegas and et al. 2007. Many Eyes: A site for visualization at internet scale. *TVGC* (2007), 1121–1128.
[18] Marcos R. et al. Vieira. 2011. On query result diversification. *ICDE*.

[19] Eugene Wu, Leilani Battle, and Samuel R. Madden. 2014. The case for data visualization management systems. *VLDB Endowment*.

[20] Cong Yu and et al. 2009. It Takes Variety to Make a World: Diversification in Recommender Systems. In *EDBT*. ACM, New York, NY, USA, 368–378.

[21] Mi Zhang and Neil Hurley. 2008. Avoiding Monotony: Improving the Diversity of Recommendation Lists. *ACM Conference on Recommender Systems*.